

Introduction to Computer Networks and the Internet

COSC 264

IP and Related Protocols

Dr. Andreas Willig

Dept. of Computer Science and Software Engineering
University of Canterbury, Christchurch

UoC, 2019

Outline

- 1 IPv4
 - Packet Format
 - IP Addressing
 - IP Forwarding and Routing
 - Fragmentation and Reassembly
- 2 IP Helper Protocols
 - ARP
 - ICMP

About This Module

- Goals of this Module:
 - Get a first idea of the Internet
 - Get to know the IP protocol and important support protocols
- Useful references:
 - The “bible” on TCP/IP: [9] (old, but still great!)
 - Other references: [4], [8, Part V]
 - Internet protocols are published as **requests-for-comment** (RFC) by the Internet Engineering Task Force (IETF), you can access them via: <http://www.ietf.org/rfc.html>
- Most of these slides are based on [9]

The Internet

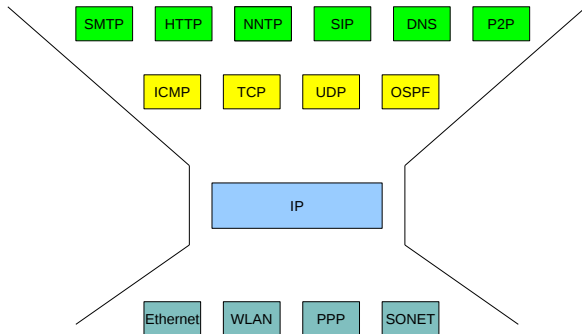
- The Internet is a packet-switched network
- It is a **network of networks**:
 - It consists of many different networks, connected by routers
 - The networks or links can be of any technology:
 - Ethernet
 - Optical point-to-point links
 - Wireless LAN
 - ...
 - Carrier pigeons (RFC 1149)
- It is really large:
 - The Internet Systems Consortium estimates ≈ 1.033 billion stations (called hosts) as of July 2015
 - See <https://www.isc.org/network/survey>
- It has a fairly complex topology [1]

The Internet (2)

- The end-to-end principle [7]:
 - Perform intelligent functions in hosts, not in routers
 - For example:
 - Routers know how to deliver packets
 - All functions making this delivery **reliable** are performed in the end host, e.g. by the TCP protocol
 - There is no network-layer mechanism for reliable delivery
 - **Keep the routers simple!**
- Internet is standardized by the IETF, standards are called RFCs
 - IETF = Internet Engineering Task Force (www.ietf.org)
 - RFC = Request For Comment
- For the design philosophy see [3]

The Hourglass Model for the Internet Protocol Stack

File sharing, WWW, Internet Telephony,



- “Everything over IP, IP over everything”

Outline

1 IPv4

2 IP Helper Protocols

Introduction

- IP is specified in RFC 791 and many followup RFCs
- It is the network layer protocol of the Internet
- Some terminology:
 - IP packets are called **datagrams**, except if they result from fragmentation and reassembly, then they are called **fragments**
 - End stations are called **hosts**
 - IP routers are called **routers**
- IP addresses are assigned to **network interfaces**:
 - When a host has three Ethernet adapters, it has three IP addresses, one for each adapter
 - Since most hosts have only one adapter, we speak of the IP address of that host

IP Service – Best Effort

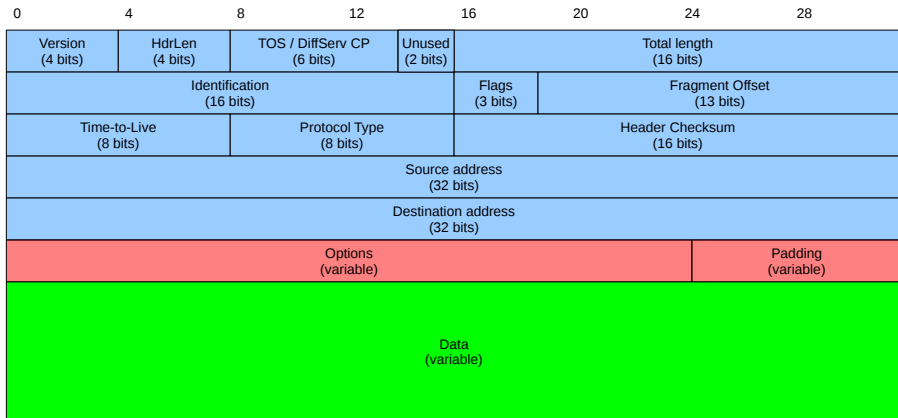
- Basic IP service is **datagram delivery**
- This service is:
 - Connectionless: no connection or shared state is set up before datagram delivery starts
 - Unacknowledged: IP does not use acknowledgements
 - Unreliable: on IP level no retransmissions are carried out
 - Unordered: IP does not guarantee in-sequence delivery [2]
- This kind of guarantee-nothing service is called **best effort**

Outline

- 1 IPv4
 - Packet Format
 - IP Addressing
 - IP Forwarding and Routing
 - Fragmentation and Reassembly

- 2 IP Helper Protocols

Packet Format



Packet Format (2)

- Where applicable (e.g. addresses), header is using **big endian** byte ordering (also called **network byte order**)
- The `Version` field specifies the version of the IP protocol running, always 4 for IPv4
- The `HdrLen` field:
 - specifies the length of IP header as number of 32-bit words
 - If the `Options` field does not use a multiple of 32 bits, a `Padding` field is used to fill up to 32 bits
 - When `HdrLen` > 5, then an `Options` field is present
- The `TOS/DSCP` field:
 - TOS = Type Of Service, DSCP = DiffServ Code Point
 - Allows to mark packets for differentiated treatment to achieve Quality-Of-Service (QoS), e.g. express priorities
 - DiffServ [5] is framework for Internet QoS, another is IntServ [10]
 - Many routers ignore the `TOS/DSCP` field

Packet Format (3)

- The `TotalLength` field:
 - Gives the total length of datagram in bytes (i.e. up to 65535)
 - Can be modified during **fragmentation and reassembly**
 - The `TotalLength` field is part of IP header, since some technologies (Ethernet!) pad up frames to achieve minimum frame size and do not reverse this
- The `Identification` field:
 - Uniquely identifies each IP payload unit accepted from higher layers for a given interface
 - In other words: it is a sequence number
 - Incremented by source host for each new IP payload
 - Routers do not touch this field
- The `Flags` field:
 - Contains two flags relevant for fragmentation and reassembly (DF, Don't Fragment, and MF, More Fragments)

Packet Format (4)

- The `FragmentOffset` field:
 - is used for fragmentation and reassembly
 - gives the offset of the current fragment within entire datagram, in multiples of eight bytes
- The `HeaderChecksum` field:
 - Is calculated over IP header only, not the data (TCP, UDP etc. have their own checksums to cover their data)
- The `Time-To-Live` field:
 - gives upper limit to number of routers a packet can traverse
 - decremented by each router, forces re-computation of checksum
 - when `TTL=1` and packet cannot be directly delivered to destination, datagram is discarded, sender is notified (ICMP message)
 - Typical initial values: 32 or 64
 - Intended usage: eliminate packets caught in a routing loop

Packet Format (5)

Protocol value	Encapsulated Protocol
0x01	ICMP
0x02	IGMP
0x04	IP-in-IP Encapsulation
0x06	TCP
0x11	UDP

- `Protocol` field indicates the higher-layer protocol that generated the payload
- This field provides **protocol multiplexing**
- In other words: provides different SAPs
- Some values shown in table

Packet Format (6)

- The `SourceAddress/DestinationAddress` fields:
 - `SrcAddr` indicates the initial sender of datagram
 - `DstAddr` indicates intended final receiver of datagram
 - Are of 32 bits width
- The `Options` field:
 - Contains header field for optional IP features
 - One example option: source routing
 - Options are rarely used, we will not consider this anymore

Outline

- 1 IPv4
 - Packet Format
 - IP Addressing
 - IP Forwarding and Routing
 - Fragmentation and Reassembly

- 2 IP Helper Protocols

IP Address Representation

- IP addresses have a width of 32 bits
- They are supposed to be worldwide unique
 - This is not really true anymore with NAT ...
- IP addresses are written in **dotted-decimal notation**, e.g.:

130.149.49.77

where decimal (!) numbers are separated by dots

- They have an internal structure:

`<network-id> <host-id>`

where:

- `<network-id>` denotes a network (e.g. an Ethernet)
- `<host-id>` refers to a host **within** this network
- The `<host-id>` must only be unique w.r.t. its network

Interlude: Routing / Forwarding Tables

- IP routers have several network interfaces or ports (different from TCP/UDP port numbers) where they receive/transmit datagrams
- In IP networks a router getting a packet on some input port looks at the `DestinationAddress` field to determine the output port
- The router consults a **forwarding table**:
 - The forwarding table lists all networks the router knows with their `<network-id>` and the output port to send the packet to in order to reach that network
 - The router performs a **table lookup** for an incoming packet, it searches the forwarding table for a matching network entry
 - Time required for table lookup depends on the number of networks stored in the table
- This is simplified, more details later!

Important Points

Important Point

A host address is tied to its location in the network, i.e. it is coupled to network topology. When a host switches to another network, it obtains another address and ongoing connections (TCP!) break – IP therefore has no direct support for mobility!!

Important Point

IP Routing is mostly concerned with networks, i.e. forwarding tables in routers mostly store `<network-id>`'s – it is the responsibility of last router on a path to deliver an IP datagram to directly connected host.

Classless Inter-Domain Routing

- Question: how many bits to allocate to `<network-id>`?
- In the early days, this number was fixed to three different values (classful addressing)
- This proved inflexible, something better was needed
- CIDR = Classless Inter-Domain Routing
- Introduced 1993, specified in RFCs 1518, 1519, mandatory
- Modern routing protocols (OSPF, RIPv2, BGP) use CIDR
- In CIDR a network is specified by two values:
 - A 32 bit network address
 - A 32 bit network mask (**netmask**)

CIDR – Netmask

- For a given 32-bit IP address the netmask specifies which bits belong to network-id and which bits belong to host-id
- The netmask consists of 32 bits, the leftmost k bits are ones, the remaining $32 - k$ bits are zeros
- Examples:

Netmask	Shorthand
11111111.11110000.00000000.00000000	/12
11111111.11111111.00000000.00000000	/16
11111111.11111111.11100000.00000000	/19
11111111.11111111.11111110.00000000	/23

- To fully specify network, give network address and netmask, e.g.:

192.168.40.0/21

- The rightmost $32 - k$ bits of a network address $a.b.c.d/k$ are zero

CIDR – Netmask (2)

- Example: given host address 192.168.40.3 and netmask /24, the hosts network address is computed as:

	11000000.10101000.00101000.00000011	192.168.40.3
AND	11111111.11111111.11111111.00000000	/24
<hr/>		
	11000000.10101000.00101000.00000000	192.168.40.0

- The same example, now with netmask /21:

	11000000.10101000.00101000.00000011	192.168.40.3
AND	11111111.11111111.11111000.00000000	/21
<hr/>		
	11000000.10101000.00101000.00000000	192.168.40.0

- In both examples the network addresses are the same, but the networks are of different size
- Hence, to distinguish both networks you need to specify both network address and bitmask, network address alone is insufficient

CIDR – Netmask (3)

- In a network $a.b.c.d/k$ there are two “special host addresses”:
 - The host address $000 \dots 00$ (with $32 - k$ zeros in total) is part of the network id, signifying that we refer to the network as a whole
 - The host address $111 \dots 11$ (with $32 - k$ ones in total) is the broadcast address of this network

All the other host addresses can be assigned to individual hosts

- Example: In the network $192.168.40.64/28$ there are 14 addresses available:
 - The netmask leaves four bits for the host-id, i.e. 16 values
 - The value 0000 is part of the network-id
 - The value 1111 is the broadcast address for this network

Reserved IP address blocks

Address Block	Current Usage
10.0.0.0/8	Private-use IP networks
127.0.0.0/8	Host loopback network
169.254.0.0/16	Link-local for point-to-point links (e.g. dialup)
172.16.0.0/12	Private-use IP networks
192.168.0.0/16	Private-use IP networks

(from: [6], there are more than shown here)

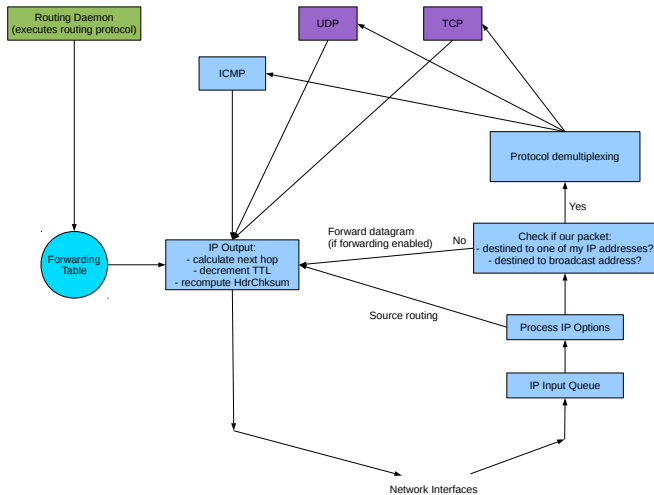
- Private-use IP addresses are often used for broadband clients or by NAT boxes
- The “traditional” loopback address of a host is 127.0.0.1, but any address from 127.0.0.0/8 network serves same purpose
- Packets with private addresses are not routed in the public internet, only within the provider network

Outline

- 1 IPv4
 - Packet Format
 - IP Addressing
 - IP Forwarding and Routing
 - Fragmentation and Reassembly

- 2 IP Helper Protocols

Simplified Packet Processing



Simplified Packet Processing (2)

- Packet processing chain is followed in routers and hosts
- Incoming packets are checked for correctness and stored in IP input queue – correctness includes:
 - right value in IP version field
 - correct IP header checksum
- Next, packet options are processed (rarely used)
- Next, it is checked if packet is destined to **this** host / router or to the broadcast address of any network this host / router is directly attached to
- If so, protocol demultiplexing is carried out
 - The `Protocol` field in IP header is checked for its value
 - Packet payload is delivered to the software entity implementing the indicated higher-layer protocol
 - Packet is not processed any further!

Simplified Packet Processing (3)

- If packet is not destined to this host/router or broadcast address:
 - If packet forwarding is not enabled, the packet is dropped
 - Otherwise:
 - Check if packet is destined to a directly reachable station (e.g. on a directly attached Ethernet) – if so, deliver packet directly
 - If packet is not destined to directly reachable station, consult forwarding table to determine next hop / outgoing interface
 - Decrement TTL value, drop packet when it reaches zero
 - Recompute packet header checksum (why?)
 - Hand packet over to outgoing interface
- Forwarding table is maintained by a **routing daemon**, i.e. a process executing a routing protocol
- Note that datagrams to be routed can come from local applications or from other hosts via IP input queue
- Linux commands to inspect / modify forwarding table:
 - `netstat`
 - `route`

Forwarding Table Contents (First Approximation)

- Each entry in the forwarding table contains:
 - Destination IP address, which can be either:
 - a full host address (i.e. non-zero host-id)
 - a network address, with netmask
 - depending on the value of a flag
 - Information about next hop, either:
 - IP address of next-hop router (must be directly reachable)
 - IP address of directly-connected network (network address)
 - Flags:
 - A flag telling whether destination IP is host or network
 - A flag telling whether next hop is a router or directly attached network
 - Specification of outgoing interface

Forwarding (First Approximation)

- From forwarding table structure it is clear that a host / router does not know the full path, but only next hop
- Forwarding table lookup for a packet with destination IP address `dst` proceeds in three stages (Caveat: reality is different):
 - First look for an entry that is a full-host address matching `dst` – if found, send packet to indicated next hop / outgoing interface and stop processing
 - This is not used very often
 - Next look for an entry that is a network address matching `dst` – if found, send packet to indicated next hop / outgoing interface and stop processing
 - Finally look for special **default** entry – if found, send packet to indicated next hop (the **default router**) and stop processing
 - Otherwise drop packet, send ICMP message back to original sender of datagram

Forwarding – Address Matching

- **Question:** how to check whether a destination address dst matches a forwarding table entry for network $a.b.c.d/k$?

Forwarding – Address Matching (2)

- **Answer:** They match when

$$(dst \text{ AND } < /k - netmask >) == (a.b.c.d \text{ AND } < /k - netmask >)$$

- **Example:** We are given the following forwarding table:

Destination Network/Netmask	Outgoing interface
130.1.0.0 / 16	eth0
141.5.6.0 / 24	eth1

- **Question:** We are given two packets with destination addresses 130.1.9.5, and 166.42.17.12, respectively. Which decisions does the router make?
- **Question:** and what happens if a default route is added to the forwarding table?

Forwarding Tables in Hosts

- Most end hosts leverage the **default route** mechanism:
 - An end host can differentiate between packets to local destinations and to all other destinations
 - **Question:** suppose an end host has address 130.149.49.77 and is part of a /24 network – how does it check whether a destination address a.b.c.d belongs to another host in the same network?
 - Packets to local destinations are delivered directly (see discussion of ARP for how to do this in an Ethernet)
 - Packets to all other destinations are sent to default router
- Therefore, forwarding tables in end hosts can be made out of just two (or more generally: a few) entries:
 - One entry for each network it is directly attached to
 - The default route
- The default route must be configured (typically done by DHCP)

Forwarding Tables in Routers

- Most routers at the “border” of the Internet only have forwarding table entries for a subset of all networks attached to the Internet (likely other routers belonging to the same owner), for all other networks they rely on default routers
- Some routers in the core:
 - do not have a default router
 - are the default routers of other routers
 - must know (almost) all the Internet networks

Outline

- 1 IPv4
 - Packet Format
 - IP Addressing
 - IP Forwarding and Routing
 - Fragmentation and Reassembly

- 2 IP Helper Protocols

On the Choice of Packet Size

- The link-layer technologies underlying IP offer many different maximally allowed packet sizes, e.g.:
 - Ethernet: 1500 bytes
 - Gigabit Ethernet: 9000 bytes
 - IEEE 802.11 WLAN: 2312 bytes
 - ISDN: 576 bytes
- This maximum size is also known as **maximum transmission unit** (MTU)
- Higher-layer protocols (TCP, UDP) and applications should not be required to know these maximal sizes:
 - One reason: “software hygiene”, separation of concerns
 - Another reason: it is not well defined:
 - Different packets of the same flow can take different routes
 - A packet can use different technologies while on travel
 - Even if all packets go the same route, this route can change due to link failures / restores

Fragmentation and Reassembly

- IP **hides** this from upper layers, offers own maximum message length of 65515 bytes to higher layers
 - $65515 = 65535 - 20$, 20 bytes is minimum size of IP header
- To cope with smaller MTUs:
 - Sender IP instance partitions message into **fragments**
 - Each fragment is transmitted individually as a full IP packet, with header information specifying that this is a fragment and giving the position of fragment in whole message
 - Each IP datagram carrying a fragment must not be larger than MTU of outgoing link
 - IP instance at destination buffers received fragments, re-assembles message and delivers to higher layers

Question

Would it be useful to have intermediate IP routers perform reassembly?

Fragmentation and Reassembly (2)

- In addition, every intermediate router can:
 - fragment a full message
 - further fragment a fragmentwhen necessary for transmission on next hop
- When the destination receives the first fragment, it:
 - Allocates buffer large enough for whole message
 - Starts a timer
- When all fragments arrive before timer expiration:
 - Timer is canceled
 - Re-assembled packet is handed over to higher layers
 - Buffer is de-allocated
- When timer expires before all fragments have arrived:
 - The already received fragments are dropped, buffer is freed
 - ICMP message (type 11, code 1) is sent to source host

Some Details

- Every message handed over to IP from higher layers has its own identifier
 - See `Identification` field in IP header
- All fragment datagrams belonging to same message have:
 - A full IP header
 - The same `Identification` field
 - A `TotalLength` field reflecting the fragment size
 - Different values for `FragmentOffset` field (reflecting the start of the present fragment within the whole message):
 - `FragmentOffset` specifies offset in multiples of 64 bits
 - The `MF` (more-fragments) bit set, except for the last fragment, which has non-zero `FragmentOffset`

Question

With this setup: how much buffer space shall the receiver allocate when collecting fragments?

Some Details: The DF bit

- By setting the DF (don't fragment) bit in the IP header a source node **forbids** fragmentation by intermediate routers
- When a router receives a datagram with DF set, it:
 - Checks whether outgoing link for this packet has an MTU large enough to transmit the packet
 - If so, the packet is transmitted onto next hop
 - If not, the router drops the datagram and returns an ICMP message to original IP source
 - ICMP with type 2 ("destination unreachable") and code 4 ("fragmentation required, but DF set")

Some Details: The DF bit (2)

Question

How could you use this for the sender to determine the **path MTU**, defined as the smallest MTU of all links along a path between source and destination?

Fragmentation and Reassembly – Discussion

- Fragmentation/Reassembly creates significant overhead:
 - Several datagrams per message, each having full IP header
 - Reassembly adds significant complexity to receiver
 - Upon loss of single fragment the whole message is possibly re-transmitted by higher layers (TCP!)

Outline

1 IPv4

2 IP Helper Protocols

Outline

1 IPv4

2 IP Helper Protocols

- ARP
- ICMP

Address Resolution Protocol – ARP

- IP addresses only have a meaning to IP and higher layers
- In an Ethernet, stations have own 48-bit MAC addresses
- Recall that IP datagrams are encapsulated into Ethernet frames
- An Ethernet station picks up a packet only if the destination MAC address matches its own MAC address (ignoring broadcast / multicast), IP addresses and other packet contents are ignored
- An IP address is assigned to an Ethernet adapter

Important Question

How do other stations know to which MAC address a given IP address refers?

Address Resolution Protocol – ARP (2)

- ARP determines MAC address for given IP address
- ARP is specified in RFC 826
- ARP is not restricted to Ethernet, but in general is geared towards LANs with broadcast capabilities
- ARP is **dynamic**:
 - The MAC address for a given IP address does not need to be statically configured, ARP allows to determine this on-the-fly
 - Advantage: nodes can be moved or equipped with new MAC adapters without any re-configuration
 - Disadvantage: a separate protocol is needed, bringing additional complexity and requiring a little bandwidth
- There is also a protocol that lets stations find an IP address for a given MAC address, this is called RARP (Reverse ARP)

Basic Operation of ARP

- Suppose that:
 - We have two stations *A* and *B* attached to the same Ethernet, having the following addresses:

	Station A	Station B
MAC	11:11:11:11:11:11	22:22:22:22:22:22
IP	130.149.49.11	130.149.49.22

- Both *A* and *B* are in the same IP network 130.149.49.00/24, which is an Ethernet network
 - Station *A* wishes to send an IP packet to address 130.149.49.22 and does not yet have any information about the corresponding MAC address
- Each station maintains an **ARP Cache**, which stores the mappings from IP to MAC addresses that the station currently knows about

Basic Operation of ARP (2)

- Station *A* **broadcasts** an **ARP-request** message (displayed in Wireshark as `arp who-has`), indicating:
 - *A*'s own IP and MAC address
 - *B*'s IP address

Broadcasting means: packet is sent to **Ethernet broadcast address**, Ethernet frame has value `0x0806` in the length/type field

- Any host *C* having an IP address other than `130.149.49.22` simply drops the ARP-request packet
- Upon receiving the ARP request, host *B* (with IP address `130.149.49.22`) performs the following actions:
 - It stores a binding between *A*'s IP and MAC address in its own ARP cache
 - It responds with an **ARP-reply packet** that includes:
 - *B*'s MAC and IP address
 - *A*'s MAC and IP address

ARP reply is unicast to *A*'s MAC addr. (Why no broadcast?)

Basic Operation of ARP (3)

- Upon receiving ARP response from *B*, station *A* stores a binding between *B*'s IP and MAC address in its ARP cache
- This procedure is called **address resolution**
- ARP makes no retransmissions when ARP request not answered
- If a station wants to send an IP packet to a local destination with address `xx.xx.xx.xx`, it:
 - first checks the ARP cache whether a binding for `xx.xx.xx.xx` can be found
 - If so, the packet is encapsulated in an Ethernet frame and directed to the MAC address found in the ARP cache
 - Otherwise, the address resolution procedure is started and the packet is sent when the result is available


The ARP Cache

- The entries in an ARP cache are soft-state, entries are typically removed 20 minutes after their creation
 - Why?
 - Some implementations restart the timer after each reference to an ARP cache entry
- Under Linux you can inspect your ARP cache with the command:

```
/usr/sbin/arp -a
```

The path to the `arp` command can vary between systems

The ARP Frame Format



(See [9, Sect. 4])

- **HardType** determines the type of MAC addresses used, `0x0001` for Ethernet 48-bit addresses
- **ProtType** determines the higher-layer protocol for which address resolution needs to be done, value `0x0800` for IP
- **HardSize** and **ProtSize** specify the size (in bytes) of the hardware and protocol addresses – they are 6 and 4 for Ethernet and IP
- **Op** distinguishes between ARP-request and ARP-reply, and some other types (RARP is covered as well)
- The remaining four fields are the mentioned address fields

Outline

1 IPv4


2 IP Helper Protocols

- ARP
- ICMP

Introduction

- ICMP = Internet Control Message Protocol
- Specified in RFC 792
- This protocol:
 - Accompanies the IP protocol by allowing routers or destination hosts to inform sender about “unusual” situations, including:
 - There is no route to the destination
 - Destination host is not reachable
 - Fragmentation required but DF set
 - Operates “on top” of IP, i.e. ICMP messages are encapsulated into regular IP datagrams
 - Does not add additional mechanisms (like error control) to IP
 - Does not force any host/router to generate ICMP messages
 - IP sending host **must not rely** on ICMP messages
- These days ICMP messages are often filtered out by firewalls

Message Format



- `type` and `code` specify actual ICMP message type and sub-type
- `checksum` covers ICMP header and data, with `checksum` assumed as zero
- Data depends on `type/code` combination, but often includes the first few bytes of the offending IP datagram (including its header)

Some type/code Combinations

type	code	Meaning
0	0	Echo reply
3	0	Destination network unreachable
3	1	Destination host unreachable
3	2	Destination protocol unreachable
3	3	Destination port unreachable
3	4	Fragmentation required, but DF bit set
3	6	Destination network unknown
3	7	Destination host unknown
4	0	Source quench (Congestion control)
8	0	Echo request
11	0	TTL expired in transit
11	1	Fragment reassembly time exceeded

- There are many more, e.g. for router advertisements, information about malformed IP packets, etc.
- It is implementation-dependent, which ICMP messages are generated
- ICMP messages are often suppressed by firewalls, otherwise too much information about internal network structures could be revealed

Some type/code Combinations (2)

- Source-quench (type=4, code=0):
 - generated by IP router when it drops a packet due to congestion
 - Intention is to let source host throttle its rate
- TTL expiration (type=11, code=0):
 - generated by IP router when it drops a packet because its TTL value reached zero
 - **Question:** this is used by `traceroute`. How?
- Fragment reassembly timeout (type=11, code=1):
 - Generated by destination when not all fragments of a message have been received within timeout
 - Used to invite higher-layer protocol at sending host to re-transmit the message (with all the fragments)
 - IP itself does not perform any retransmission!

Some type/code Combinations (3)

- The “destination-unreachable” messages (`type=3`):
 - `code=0` (destination network unreachable) and `code=1` (destination host unreachable): generated when:
 - `code=0`: router does not have a matching entry (and no default entry) for non-directly connected destination address in its forwarding table
 - `code=1`: router could not deliver datagram to directly connected host (e.g. no ARP response)
 - `code=2` (protocol unreachable): IP datagram refers to non-existent higher-layer protocol in destination (`protocol-type` field)
 - `code=3` (port unreachable): used with TCP / UDP
 - In these messages first 32 bits of the variable ICMP message part are 0, following bytes contain IP header and first few bytes of offending IP datagram

- [1] David Alderson, Lun Li, Walter Willinger, and John C. Doyle.
Understanding Internet Topology: Principles, Models and Validation.
IEEE/ACM Transactions on Networking, 13(6):1205–1218, December 2005.
- [2] Jon C. R. Bennett, Craig Partridge, and Nicholas Shectman.
Packet Reordering is Not Pathological Network Behaviour.
IEEE/ACM Transactions on Networking, 7(6):789–798, December 1999.
- [3] David D. Clark.
The design philosophy of the darpa internet protocols.
ACM Computer Communication Review, 18(4):106–114, August 1988.
- [4] Douglas E. Comer.
Internetworking with TCP/IP – Principles, Protocols and Architecture, volume 1.
Prentice Hall, Englewood Cliffs, New Jersey, third edition, 1995.
- [5] Kalevi Kilkki.
Differentiated Services for the Internet.
Macmillan Technical Publishing, Indianapolis, 1999.
- [6] Deepankar Medhi and Karthikeyan Ramasamy.
Network Routing – Algorithms, Protocols, and Architectures.
Morgan Kaufmann, San Francisco, California, 2007.
- [7] Jerome H. Saltzer, David P. Reed, and David D. Clark.
End-to-end arguments in system design.
ACM Transactions on Computer Systems, 2(4):277–288, November 1984.

- [8] William Stallings.
Data and Computer Communications.
Prentice Hall, Englewood Cliffs, New Jersey, fourth edition, 2006.
- [9] W. Richard Stevens.
TCP/IP Illustrated Volume 1 - The Protocols.
Addison-Wesley, Boston, Massachusetts, 1995.
- [10] Paul P. White and Jon Crowcroft.
The integrated services in the internet: State of the art.
Proceedings of the IEEE, 85(12):1934–1946, December 1997.