

COSC264

Introduction to Computer Networks and the Internet

Introduction to Routing- Link State Routing

Dr. Barry Wu

Wireless Research Centre


University of Canterbury

barry.wu@canterbury.ac.nz


A quick review

- Layer approach and services
- Hierarchical routing and Autonomous System (AS)
- Routing vs forwarding
- Classification of routing algorithms

Hierarchical routing in the Internet



Routing and Forwarding



Outline – today


- Network layer overview
- Routing overview
- **Link-state routing (Dijkstra's algorithm)**
- Distance-vector routing (Bellman-Ford)
- Summary

Routing Algorithms and Routing Protocols

Intra-AS Routing

Routing Protocols	Routing Algorithms
RIP	Bellman-Ford (Distance-vector) Algorithm
OSPF	Dijkstra's Algorithm
BGP	Bellman-Ford (Distance-vector) Algorithm

Inter-AS Routing


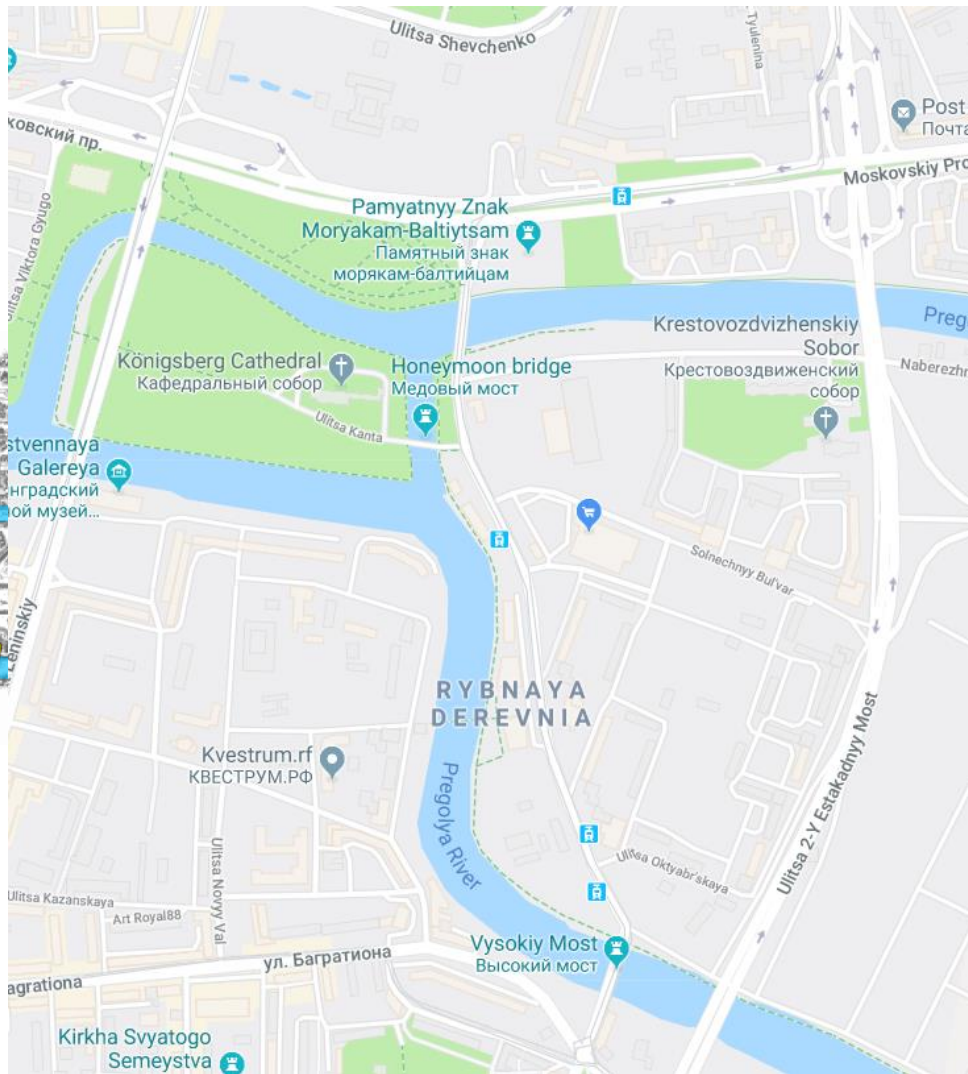



The Internet routing protocols (RIP, OSPF, and BGP) are *load-insensitive*.

Euler and Graph Theory

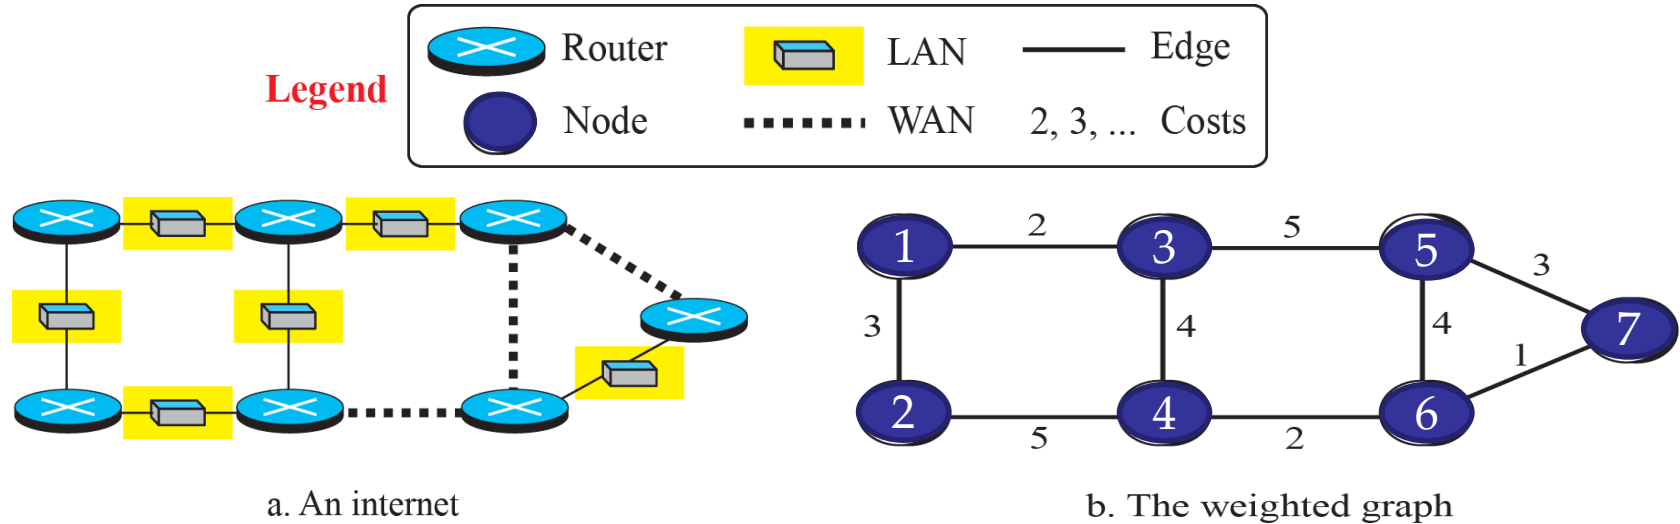
- Seven Bridges of Königsberg. 1783

- [Wiki](#)



Modeling a network

■ A network and its graphical representation



A graph G as an ordered pair $G = (V, E)$ where


- V is a set of nodes (vertices) for routers; nodes are vertices $v_i \in V$
 - e.g., $V = \{1, 2, 3, \dots, N\}$
- E is a set of edges (links); $e_{ij} = (v_i, v_j) \in E$
 - $E \subset V \times V$; $E = \{(1, 2), (1, 3), (2, 4), (3, 4), (3, 5), (4, 6), (5, 6), (5, 7), (6, 7)\}$
 - v_i and v_j are neighbors
 - Edge weights are costs

Modeling a network (2)

- Modeled as a graph

- Routers \Rightarrow nodes
- Link \Rightarrow edges

Another example network graph





- Edge labels (called metrics) can be interpreted differently

- as costs, e.g., delay, monetary transmission costs, geographical distance
- as available resources, e.g., number of available phone trunks, current available capacity given the set of flows that already use this link

Routing algorithms

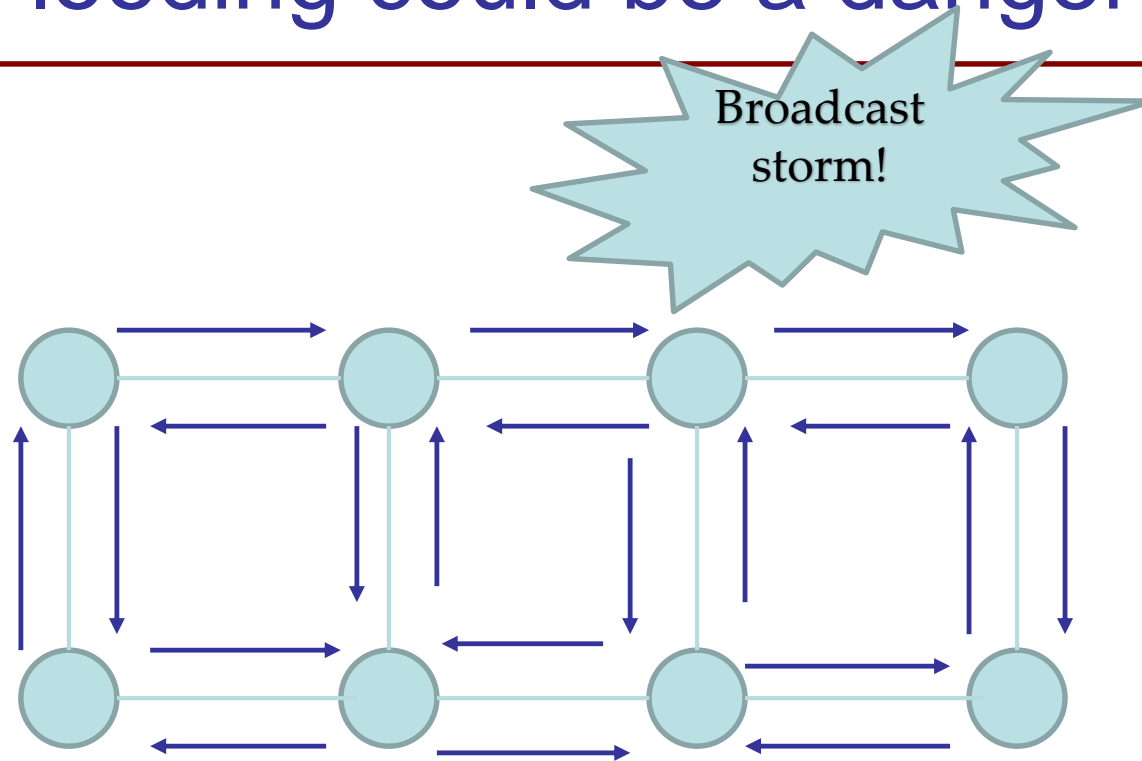
- To find least cost path
 - **Shortest path** if all link costs equal (measures hops)



Link State Routing

- Each router has **complete network picture**
 - Topology, Link costs
- How does each router get the global state?
 - Each router reliably **floods** information about its neighbors to every other router; authentication;
 - All routers have consistent information;


Flooding could be a danger!




There are sophisticated algorithms doing the broadcasting job! (Controlled flooding, spanning-tree broadcast; refer to 4.7 of [KR3].)

Link State: Control Traffic

- Each node floods its local information
- Each node ends up knowing the *entire* network topology




Link State: Node State



Link State Routing

- Each router independently calculates the least-cost path from itself to every other router
 - Using Dijkstra's Algorithm;
 - Generates a forwarding table for every destination;



Dest.	Next-hop
u	x
v	y
...	...

Dijkstra's Algorithm

- INPUT:
 - Network topology (graph), with link costs


- OUTPUT:
 - Least cost paths from one node to all other nodes

Dijkstra's Algorithm

- **S**: nodes whose least-cost path already known
 - Initially, $\mathbf{S} = \{u\}$ where u is the source node
 - Add one node to **S** in each iteration
- **D(v)**: current cost of path from source to node v
 - Initially, $\mathbf{D}(v) = \mathbf{c}(u, v)$ for all nodes v adjacent to u
 - ... and $\mathbf{D}(v) = \infty$ for all other nodes v
 - Continually update **D(v)** as shorter paths are learned
- $p(v)$: predecessor node along path from source to v , that is next to v


Dijkstra's Algorithm

```
1  Initialization:
2   $S = \{u\}$  /*  $u$  is the source */
3  for all nodes  $v$ 
4    if  $v$  is adjacent to  $u$  {
5      then  $D(v) = c(u, v)$  /* cost of neighbor known */
6      else  $D(v) = \infty$  /* cost of others unknown */
7
8  Loop
9    find  $w$  not in  $S$  with the smallest  $D(w)$ 
10   add  $w$  to  $S$ 
11   update  $D(v)$  for all  $v$  adjacent to  $w$  and not in  $S$ :
12      $D(v) = \min\{D(v), D(w) + c(w, v)\}$ 
13     /* new cost to  $v$  is either old cost to  $v$  or known
        shortest path cost to  $w$  plus cost from  $w$  to  $v$  */
14 until all nodes in  $S$ 
```




Dijkstra's Algorithm with another example

- $c(i,j)$: link cost from node i to j ; cost infinite if not direct neighbors; ≥ 0
- $D(v)$: current value of cost of path from source to destination v
- $p(v)$: predecessor node along path from source to v , that is next to v
- S : set of nodes whose least cost path definitively known



Example: Dijkstra's Algorithm

Step	start S	D(B),p(B)	D(C),p(C)	D(D),p(D)	D(E),p(E)	D(F),p(F)
→ 0	A	2,A	5,A	1,A	∞	∞
1						
2						
3						
4						
5						




```

1  Initialization:
2  S = {A};
3  for all nodes v
4    if v is adjacent to A
5      then D(v) = c(A,v);
6      else D(v) =  $\infty$ ;
...
  
```

Example: Dijkstra's Algorithm

Step	start S	D(B),p(B)	D(C),p(C)	D(D),p(D)	D(E),p(E)	D(F),p(F)
0	A	2,A	5,A	1,A	∞	∞
1						
2						
3						
4						
5						




```

...
8  Loop
9  find w not in S s.t.  $D(w)$  is a minimum;
10 add w to S;
11 update  $D(v)$  for all v adjacent
    to w and not in S:
12 If  $D(w) + c(w,v) < D(v)$  then  $D(v) =$ 
     $D(w) + c(w,v)$ ;  $p(v) = w$ ;
13 until all nodes in S;
    
```

Example: Dijkstra's Algorithm

Step	start S	D(B),p(B)	D(C),p(C)	D(D),p(D)	D(E),p(E)	D(F),p(F)
0	A	2,A	5,A	1,A	∞	∞
1	AD					
2						
3						
4						
5						




```

...
8  Loop
9   find w not in S s.t. D(w) is a minimum;
10  add w to S;
11  update D(v) for all v adjacent
    to w and not in S:
12  If  $D(w) + c(w,v) < D(v)$  then
     $D(v) = D(w) + c(w,v)$ ;  $p(v) = w$ ;
13  until all nodes in S;
    
```

Example: Dijkstra's Algorithm

Step	start S	D(B),p(B)	D(C),p(C)	D(D),p(D)	D(E),p(E)	D(F),p(F)
0	A	2,A	5,A	1,A	∞	∞
1	AD		4,D		2,D	
2						
3						
4						
5						




```

...
8  Loop
9  find w not in S s.t. D(w) is a minimum;
10 add w to S;
11 update D(v) for all v adjacent
    to w and not in S:
12 If  $D(w) + c(w,v) < D(v)$  then
     $D(v) = D(w) + c(w,v)$ ;  $p(v) = w$ ;
13 until all nodes in S;
    
```

Example: Dijkstra's Algorithm

Step	start S	D(B),p(B)	D(C),p(C)	D(D),p(D)	D(E),p(E)	D(F),p(F)
0	A	2,A	5,A	1,A	∞	∞
1	AD		4,D		2,D	
2	ADE		3,E			4,E
3						
4						
5						




```

...
8  Loop
9   find w not in S s.t. D(w) is a minimum;
10  add w to S;
11  update D(v) for all v adjacent
    to w and not in S:
12  If D(w) + c(w,v) < D(v) then
    D(v) = D(w) + c(w,v); p(v) = w;
13  until all nodes in S;
    
```


Example: Dijkstra's Algorithm

Step	start S	D(B),p(B)	D(C),p(C)	D(D),p(D)	D(E),p(E)	D(F),p(F)
0	A	2,A	5,A	1,A	∞	∞
1	AD		4,D		2,D	
2	ADE		3,E			4,E
→ 3	ADEB					
4						
5						




```

...
8  Loop
9   find w not in S s.t.  $D(w)$  is a minimum;
10  add w to S;
11  update  $D(v)$  for all v adjacent
    to w and not in S:
12  If  $D(w) + c(w,v) < D(v)$  then
     $D(v) = D(w) + c(w,v)$ ;  $p(v) = w$ ;
13  until all nodes in S;
    
```

Example: Dijkstra's Algorithm

Step	start S	D(B),p(B)	D(C),p(C)	D(D),p(D)	D(E),p(E)	D(F),p(F)
0	A	2,A	5,A	1,A	∞	∞
1	AD		4,D		2,D	
2	ADE		3,E			4,E
3	ADEB					
4	ADEBC					
5						




```

...
8  Loop
9   find w not in S s.t. D(w) is a minimum;
10  add w to S;
11  update D(v) for all v adjacent
    to w and not in S:
12  If D(w) + c(w,v) < D(v) then
    D(v) = D(w) + c(w,v); p(v) = w;
13  until all nodes in S;
    
```

Example: Dijkstra's Algorithm

Step	start S	D(B),p(B)	D(C),p(C)	D(D),p(D)	D(E),p(E)	D(F),p(F)
0	A	2,A	5,A	1,A	∞	∞
1	AD		4,D		2,D	
2	ADE		3,E			4,E
3	ADEB					
4	ADEBC					
5	ADEBCF					




```

...
8  Loop
9   find w not in S s.t. D(w) is a minimum;
10  add w to S;
11  update D(v) for all v adjacent
    to w and not in S:
12  If D(w) + c(w,v) < D(v) then
    D(v) = D(w) + c(w,v); p(v) = w;
13  until all nodes in S;
    
```

Example: Dijkstra's Algorithm


Step	start S	D(B),p(B)	D(C),p(C)	D(D),p(D)	D(E),p(E)	D(F),p(F)
0	A	2,A	5,A	1,A	∞	∞
1	AD		4,D		2,D	
2	ADE		3,E			4,E
3	ADEB					
4	ADEBC					
5	ADEBCF					



To determine path $A \rightarrow C$ (say),
work backward from C via $p(v)$

The Forwarding Table

- Running Dijkstra at node *A* gives the shortest path from *A* to all destinations
- We then construct the *forwarding* table




Subnet address
(CIDRised) and
interface

Destination	Link
B	(A,B)
C	(A,D)
D	(A,D)
E	(A,D)
F	(A,D)

Dijkstra's algorithm, discussion


Algorithm complexity: n nodes

- each iteration: need to check all nodes, w , not in N
- $n(n-1)/2$ comparisons: $O(n^2)$
- more efficient implementations possible: $O(n \log n)$
 - Using a min-heap;
 - we can find out the node with min cost in $O(\log n)$;
 - Total cost = $O(\log(n-1) + \log(n-2) + \dots + \log 1) = O(\log(n!))$
 - $= O(n \log n)$ (using *Stirling's approximation*).




Oscillation with link-state routing


Today's Internet routing algorithms are load-*insensitive*!




a. Initial routing



b. x,y detect better path to w, clockwise



c. x, y, z detect better path to w, counterclockwise




d. x,y,z detect better path to w, clockwise

Summary: Link-State Routing

- Each router broadcasts the link state
 - To give every router a complete view of the graph
- Each router runs Dijkstra's algorithm
 - Compute least-cost paths, then construct forwarding table

Exercise


Example net. 1 (Fig. 19.1 W. Stallings' book)



Dijkstra's algorithm (for node 1)

step	S	D(2), p(2)	D(3), p(3)	D(4), p(4)	D(5), p(5)	D(6), p(6)
0	{1}	2,1	5,1	1,1	inf	inf
1	{1,4}	2,1	4, 4	1,1	2,4	inf
2	{1,4,2}	2,1	4, 4	1,1	2,4	inf
3	{1,4,2,5}	2,1	3,5	1,1	2,4	4,5
4	{1,4,2,5,3}	2,1	3,5	1,1	2,4	4,5
5	{1,4,2,5,3 6}	2,1	3,5	1,1	2,4	4,5

Example net 2



Dijkstra's algorithm (for node A)

step	S	D(B), p(B)	D(C), p(C)	D(D), p(D)	D(E), p(E)	D(F), p(F)
0	{A}	1,A	inf	4,A	inf	inf
1	{A,B}	1,A	4,B	4,A	2,B	inf
2	{A,B,E}	1,A	3,E	3,E	2,B	6,E
3	{A,B,E,D}	1,A	3,E	3,E	2,B	6,E **
4	{A,B,E,D, C}	1,A	3,E	3,E	2,B	5, C
5	{A,B,E,D, C,F}	1,A	3,E	3,E	2,B	5, C

**6,E will be 5, C if node C is chosen first.

References

- [KR3] James F. Kurose, Keith W. Ross, *Computer networking: a top-down approach featuring the Internet*, 3rd edition.
- [PD5] Larry L. Peterson, Bruce S. Davie, *Computer networks: a systems approach*, 5th edition
- [TW5] Andrew S. Tanenbaum, David J. Wetherall, *Computer network*, 5th edition
- [LHBi]Y-D. Lin, R-H. Hwang, F. Baker, *Computer network: an open source approach*, International edition

Acknowledgements

- All slides are developed based on slides from the following two sources:
 - Dr DongSeong Kim's slides for COSC264, University of Canterbury;
 - Prof Aleksandar Kuzmanovic's lecture notes for CS340, Northwestern University,
https://users.cs.northwestern.edu/~akuzma/classes/CS340-w05/lecture_notes.htm