

COSC264

Introduction to Computer Networks and the Internet

Introduction to Routing – Distance Vector Algorithm

Dr. Barry Wu

Wireless Research Centre

University of Canterbury

barry.wu@canterbury.ac.nz

Outline – this week

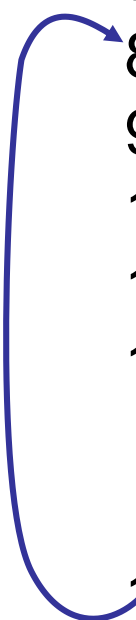
- Network layer overview
- Routing overview
- Link-state routing (Dijkstra's algorithm)
- Distance-vector routing (Bellman-Ford)
- Summary

Outline – this week

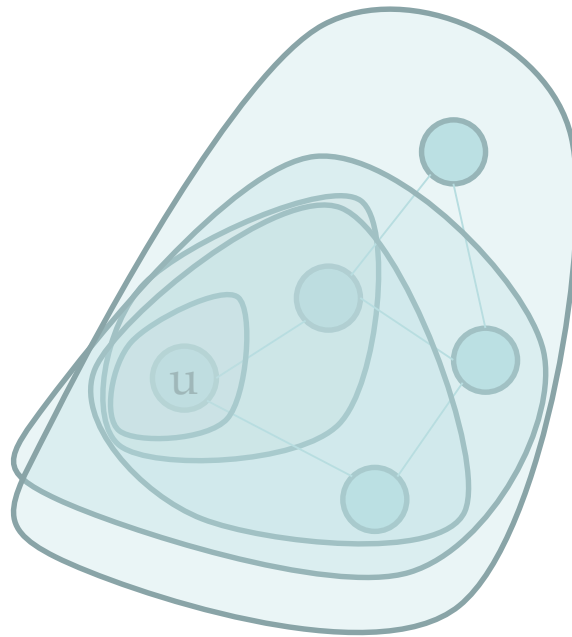
- Network layer overview
- Routing overview
- Link-state routing (Dijkstra's algorithm)
- **Distance-vector routing (Bellman-Ford)**
- **Summary**

Review: Dijkstra's Algorithm

```
1  Initialization:
2   $S = \{u\}$  /*  $u$  is the source */
3  for all nodes  $v$ 
4    if  $v$  adjacent to  $u$  {
5      then  $D(v) = c(u, v)$  /* cost of neighbor known */
6      else  $D(v) = \infty$  /* cost of others unknown */
7
8  Loop
9    find  $w$  not in  $S$  with the smallest  $D(w)$ 
10   add  $w$  to  $S$ 
11   update  $D(v)$  for all  $v$  adjacent to  $w$  and not in  $S$ :
12      $D(v) = \min\{D(v), D(w) + c(w, v)\}$ 
13     /* new cost to  $v$  is either old cost to  $v$  or known
        shortest path cost to  $w$  plus cost from  $w$  to  $v$  */
14 until all nodes in  $S$ 
```



An illustration



Outline – today

- Network layer overview
- Routing overview
- Link-state routing (Dijkstra's algorithm)
- **Distance-vector routing (Bellman-Ford)**
- **Summary**

Distance Vector Algorithm

- Distributed
- Iterative
- Asynchronous

Bellman-Ford Equation

Define


$d_x(y) :=$ cost of least-cost path from x to y

Then

$$d_x(y) = \min_v \{ c(x,v) + d_v(y) \}$$

where \min is taken over all neighbors of x

Bellman-Ford example



Node that achieves minimum is next hop in shortest path.


Clearly, $d_v(z) = 5$, $d_x(z) = 3$, $d_w(z) = 3$

B-F equation says:


$$\begin{aligned} d_u(z) &= \min \{ c(u,v) + d_v(z), \\ &\quad c(u,x) + d_x(z), \\ &\quad c(u,w) + d_w(z) \} \\ &= \min \{ 2 + 5, \\ &\quad 1 + 3, \\ &\quad 5 + 3 \} = 4 \end{aligned}$$

Distance Vector Algorithm

- Estimates:
 - $D_x(y)$ = estimate of least cost from x to y
 - Distance vector: $\mathbf{D}_x = [D_x(y): y \in N]$
- Each node x :
 - Node x knows cost to each neighbor v : $c(x,v)$
 - Node x maintains $\mathbf{D}_x = [D_x(y): y \in N]$
 - Node x also maintains its neighbors' distance vectors
 - For each neighbor v , x maintains $\mathbf{D}_v = [D_v(y): y \in N]$



An illustration



Distance vectors at node u

D_u

D_v

D_x

Distance vector algorithm

Basic idea:

- Each node periodically sends its own distance vector estimate to neighbours
- When a node x receives new DV estimate from neighbor, it updates its own DV using B-F equation:

$$D_x(y) \leftarrow \min_v \{c(x,v) + D_v(y)\} \quad \text{for each node } y \in N$$

- Amazingly, as long as all the nodes continue to exchange their distance vectors in an asynchronous fashion, the estimate $D_x(y)$ converges the actual least cost $d_x(y)$

Distance Vector Algorithm

Iterative, asynchronous:

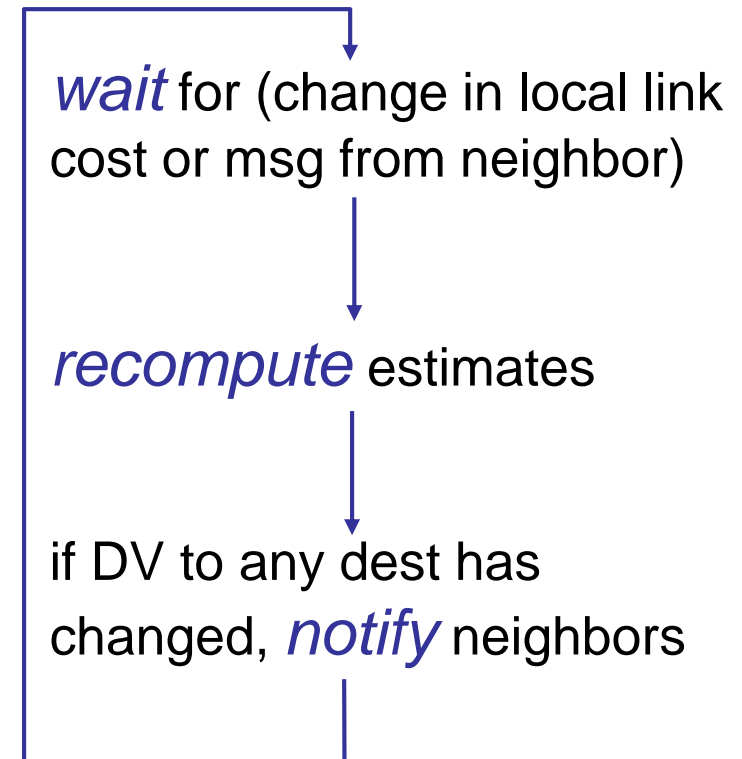
each local iteration caused by:

- local link cost change
- DV update message from neighbour

Distributed:

- each node notifies neighbours *only* when its DV changes
 - neighbours then notify their neighbours if necessary
 - The algorithm doesn't know the entire path – only knows the next hop

Each node:



Distance Vector Algorithm

At each node, x:

```
1  Initialization:
2      for all destinations y in N:
3           $D_x(y) = c(x,y)$  /*  $c(x,y) = \infty$  if y is not a neighbour */
4      for each neighbour w
5           $D_w(y) = \infty$  for all destinations y in N
6      for each neighbor w
7          send distance vector  $D_x = [D_x(y): y \text{ in } N]$  to w
8  loop
9      wait (until I see a link cost change to some neighbor w
10     or until I receive a distance vector from some neighbour w)
11     for each y in N:
12          $D_x(y) = \min_v \{c(x,v) + D_v(y)\}$  /* v is adjacent to x */
13     if  $D_x(y)$  changed for any destination y
14         send distance vector  $D_x = [D_x(y): y \text{ in } N]$  to all neighbours
15 forever
```

$$D_x(y) = \min\{c(x,y) + D_y(y), c(x,z) + D_z(y)\}$$

$$= \min\{2+0, 7+1\} = 2$$

$$D_x(z) = \min\{c(x,y) + D_y(z), c(x,z) + D_z(z)\}$$

$$= \min\{2+1, 7+0\} = 3$$

node x table

		cost to		
		x	y	z
from	x	0	2	7
	y	∞	∞	∞
	z	∞	∞	∞

cost to

		cost to		
		x	y	z
from	x	0	2	3
	y	2	0	1
	z	7	1	0

cost to

		cost to		
		x	y	z
from	x	0	2	3
	y	2	0	1
	z	3	1	0

node y table

		cost to		
		x	y	z
from	x	∞	∞	∞
	y	2	0	1
	z	∞	∞	∞

cost to

		cost to		
		x	y	z
from	x	0	2	7
	y	2	0	1
	z	7	1	0

cost to

		cost to		
		x	y	z
from	x	0	2	3
	y	2	0	1
	z	3	1	0

node z table


		cost to		
		x	y	z
from	x	∞	∞	∞
	y	∞	∞	∞
	z	7	1	0

cost to

		cost to		
		x	y	z
from	x	0	2	7
	y	2	0	1
	z	3	1	0

cost to

		cost to		
		x	y	z
from	x	0	2	3
	y	2	0	1
	z	3	1	0



time →

A hidden assumption – N (all destinations)

- Q: if all nodes exchange distance vectors with their neighbours only, can each of them know all the destination nodes (N, in the pseudocode)?




- Initially, x knows it has a path to y with cost 1; *but it does not know z*;
- y knows it has a path to x (and z) with cost 1;
- z knows it has a path to y with cost 1; *but it does not know x*;
- Then, y and x exchange distance vectors;
- x learned that there is *a new destination z* and it can reach z via y with cost 2;
- y and z exchange distance vectors;
- z learned that there is *a new destination x* and it can reach x via y with cost 2;
- Now both x and z know their destination nodes!

Distance Vector: link cost changes

Link cost changes:

- ❑ node detects local link cost change
- ❑ updates routing info, recalculates distance vector
- ❑ if DV changes, notify neighbors



“good
news
travels
fast”

At time t_0 , y detects the link-cost change ($4 \rightarrow 1$), updates its DV, and informs its neighbors.

At time t_1 , z receives the update from y and updates its table. It computes a new least cost to x ($5 \rightarrow 2$) and sends its neighbors its DV.

At time t_2 , y receives z 's update and updates its distance table. y 's least costs do not change and hence y does *not* send any message to z .

Distance Vector: link cost changes



Link cost changes:

❑ Before the link cost changes

- $D_y(x) = 4$, $D_z(x) = 5$

❑ At time t_0 , y detects the link-cost change and re-compute its dv


- $D_y(x) = \min\{c(y,x) + D_x(x), c(y,z) + D_z(x)\} = \min\{60+0, 1+5\} = 6$;
- Now we have routing loop: y sends data to z in order to get to x ; z will send data back to y in order to get to x ;
- At time t_1 , y sends its new dv to z ; after z receives y 's new dv; z can update
 $D_z(x) = \min\{c(z,y) + D_y(x), c(z,x) + D_x(x)\} = \min\{1+6, 50+0\} = 7$;
- At time t_2 , z sends its new dv to y ; similarly y can update
 $D_y(x) = \min\{c(y,x) + D_x(x), c(y,z) + D_z(x)\} = \min\{60+0, 1+7\} = 8$;
- Then $D_z(x) = \min\{c(z,y) + D_y(x), c(z,x) + D_x(x)\} = \min\{1+8, 50+0\} = 9$;
- $D_y(x) = \min\{c(y,x) + D_x(x), c(y,z) + D_z(x)\} = \min\{60+0, 1+9\} = 10$;
- ...
- $D_z(x) = \min\{c(z,y) + D_y(x), c(z,x) + D_x(x)\} = \min\{1+50, 50+0\} = 50$;
- $D_y(x) = \min\{c(y,x) + D_x(x), c(y,z) + D_z(x)\} = \min\{60+0, 1+50\} = 51$;
- $D_z(x) = \min\{c(z,y) + D_y(x), c(z,x) + D_x(x)\} = \min\{1+51, 50+0\} = 50$;



Distance Vector: link cost changes

Link cost changes:

- Before the link cost changes
 - $D_y(x) = 4$, $D_z(x) = 5$
- At time t_0 , y detects the link-cost change and re-compute its dv



- ...
- $D_z(x) = \min\{c(z,y) + D_y(x), c(z,x) + D_x(x)\} = \min\{1 + 9999, 9999 + 0\} = 9999$;
- $D_y(x) = \min\{c(y,x) + D_x(x), c(y,z) + D_z(x)\} = \min\{10000 + 0, 1 + 9999\} = 10000$;
- $D_z(x) = \min\{c(z,y) + D_y(x), c(z,x) + D_x(x)\} = \min\{1 + 10000, 9999 + 0\} = 9999$;

The process stops after z computes the cost of its path via y to be greater than 50; then it chooses the path $z \rightarrow x$ (cost is 50).

The bad news about the increase in link cost has travelled slowly!


What if $c(y,x)$ had changed from 4 to 10,000 and $c(z,x)$ had been 9,999? **Count-to-infinity** problem!

Distance Vector: link cost changes

Poisoned reverse:

□ If z routes through y to get to x :

- z tells y its (z's) distance to x is infinite (so y won't route to x via z)



□ Before the link cost changes

- $D_y(x) = 4$, $D_z(x) = 5$, but z will lie to y saying “ $D_z(x) = \infty$ ” (*poisoned reverse*)

□ At time t0, y detects the link-cost change and re-compute its dv

- $D_y(x) = \min\{c(y,x) + D_x(x), c(y,z) + D_z(x)\} = \min\{60+0, 1 + \infty\} = 60$;
- Now y sends data directly to x;
- At time t1, y sends its new dv to z; after z receives y's new dv; z can update $D_z(x) = \min\{c(z,y) + D_y(x), c(z,x) + D_x(x)\} = \min\{1+60, 50+0\} = 50$;
- At time t2, z sends its new dv to y without lying since it will not route through y; similarly y can update $D_y(x) = \min\{c(y,x) + D_x(x), c(y,z) + D_z(x)\} = \min\{60+0, 1+50\} = 51$;
- Then $D_z(x) = \min\{c(z,y) + D_y(x), c(z,x) + D_x(x)\} = \min\{1 + \infty, 50+0\} = 50$; y lies to z this time because it routes through z;

Looks like this little lie does help!

But keep lying does not *solve* the problem!

Consider y,z,w distance table entries to x only. Using poisoned reverse,

$z \rightarrow w, D_z(x) = \infty; z \rightarrow y D_z(x) = 6;$


$w \rightarrow y, D_w(x) = \infty; w \rightarrow z D_w(x) = 5;$

$y \rightarrow w, D_y(x) = 4; y \rightarrow z D_y(x) = 4;$

Then there is link-cost change (4 \rightarrow 60);

At t1, y updates its $D_y(x) = 9$ (via z); $y \rightarrow w, D_y(x) = 9; y \rightarrow z D_y(x) = \infty;$

$D_y(x) = \min\{c(y,z)+D_z(x), c(y,w) + D_w(x), c(y,x) + D_x(x)\} = \min\{3+6, 1+\infty, 60+0\} = 9$ (via z)



	t0	t1	t2	t3	t4
z	$\rightarrow w, D_z(x) = \infty;$ $\rightarrow y, D_z(x) = 6;$		No change	$\rightarrow w, D_z(x) = \infty;$ $\rightarrow y, D_z(x) = 11;$	
w	$\rightarrow y, D_w(x) = \infty;$ $\rightarrow z D_w(x) = 5;$		$\rightarrow y, D_w(x) = \infty;$ $\rightarrow z, D_w(x) = 10;$		No change
y	$\rightarrow w, D_y(x) = 4;$ $\rightarrow z, D_y(x) = 4;$	$\rightarrow w, D_y(x) = 9;$ $\rightarrow z, D_y(x) = \infty;$		No change	$\rightarrow w, D_y(x) = 14;$ $\rightarrow z D_y(x) = \infty;$

This continues y-w-z-y-w-z-y-w-z; there is routing loop (y-z, z-w, w-y) [ZL]

Comparison of LS and DV algorithms

Message complexity

- LS: with n nodes, E links, $O(nE)$ msgs sent
- DV: exchange between neighbors only
 - convergence time varies

Speed of Convergence

- LS: $O(n^2)$ algorithm requires $O(nE)$ msgs
 - may have oscillations
- DV: convergence time varies
 - may be routing loops
 - count-to-infinity problem

Robustness: what happens if router malfunctions?

LS:

- node can advertise incorrect *link* cost
- each node computes only its *own* table

DV:

- DV node can advertise incorrect *path* cost
- each node's table used by others
 - error propagate thru network

Summary

- Network layer overview
- Routing overview
- Link-state routing (Dijkstra's algorithm)
- **Distance-vector routing (Bellman-Ford)**
 - B-F equation
 - B-F algorithm
 - Count-to-infinity problem and poisoned reverse
 - LS vs DV
- **Summary**

References

- [KR3] James F. Kurose, Keith W. Ross, *Computer networking: a top-down approach featuring the Internet*, 3rd edition.
- [PD5] Larry L. Peterson, Bruce S. Davie, *Computer networks: a systems approach*, 5th edition
- [TW5] Andrew S. Tanenbaum, David J. Wetherall, *Computer network*, 5th edition
- [LHBi] Y-D. Lin, R-H. Hwang, F. Baker, *Computer network: an open source approach*, International edition
- [ZL] Lilin Zhang, CSC358 Tutorial 9, University of Toronto,
<http://www.cs.toronto.edu/~ahchinaei/teaching/2016jan/csc358/Tut09-taSlides.pdf>

Acknowledgements

- All slides are developed based on slides from the following two sources:
 - Dr DongSeong Kim's slides for COSC264, University of Canterbury;
 - Prof Aleksandar Kuzmanovic's lecture notes for CS340, Northwestern University,
https://users.cs.northwestern.edu/~akuzma/classes/CS340-w05/lecture_notes.htm