

COSC264

Introduction to Computer Networks and the Internet

# Transport Layer Protocols: UDP and TCP

Dr. Barry Wu

Wireless Research Centre

University of Canterbury

[barry.wu@canterbury.ac.nz](mailto:barry.wu@canterbury.ac.nz)

# Outline

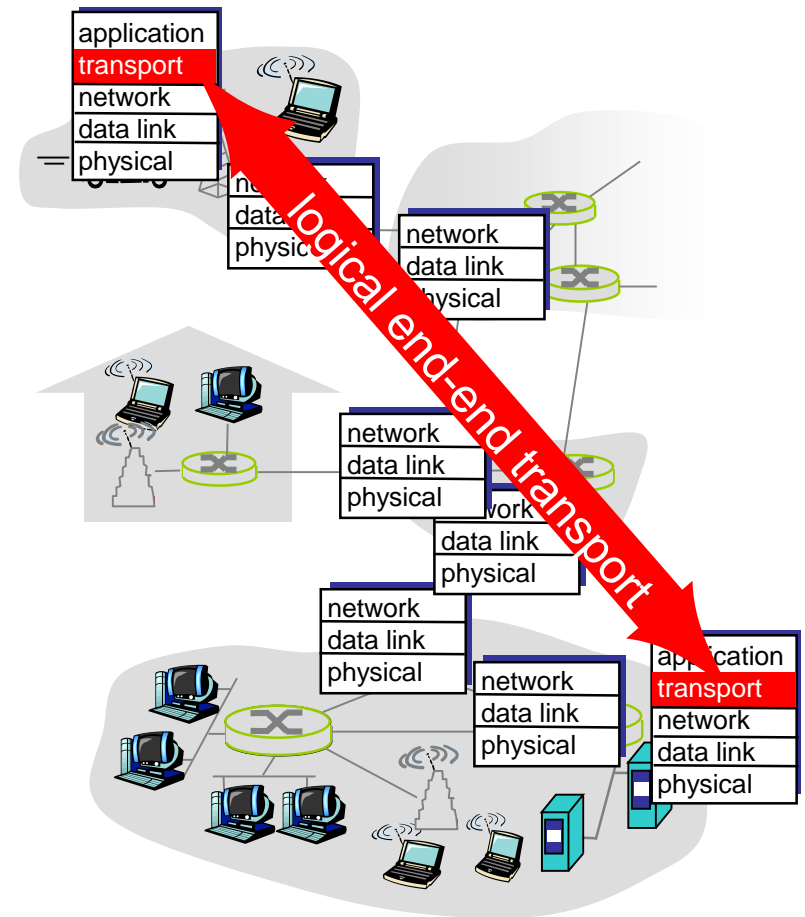
- Transport-layer services
- Multiplexing and demultiplexing
- Connectionless transport: UDP

# Role of Transport Layer

- Application layer
  - Communication for specific applications
  - e.g., HyperText Transfer Protocol (HTTP), File Transfer Protocol (FTP), Network News Transfer Protocol (NNTP)
- **Transport layer**
  - Communication between processes (e.g., socket = IP address and Port)
  - Relies on network layer and serves the application layer
  - e.g., TCP and UDP
- Network layer
  - Logical communication between nodes
  - Hides details of the link technology
  - e.g., Internet Protocol (IP)

# Transport Protocols

- Provide *logical communication* between application processes running on different hosts
- Run on end hosts
  - Sender: breaks application messages into **segments**, and passes to network layer
  - Receiver: reassembles segments into messages, passes to application layer
- Multiple transport protocol available to applications
  - Internet: TCP and UDP
    - Other: Datagram Congestion Control Protocol (DCCP), Stream Control Transmission Protocol (SCTP)



# Transport vs. network layer

- *network layer*: logical communication between hosts
- *transport layer*: logical communication between processes
  - relies on, enhances, network layer services

## Household analogy:

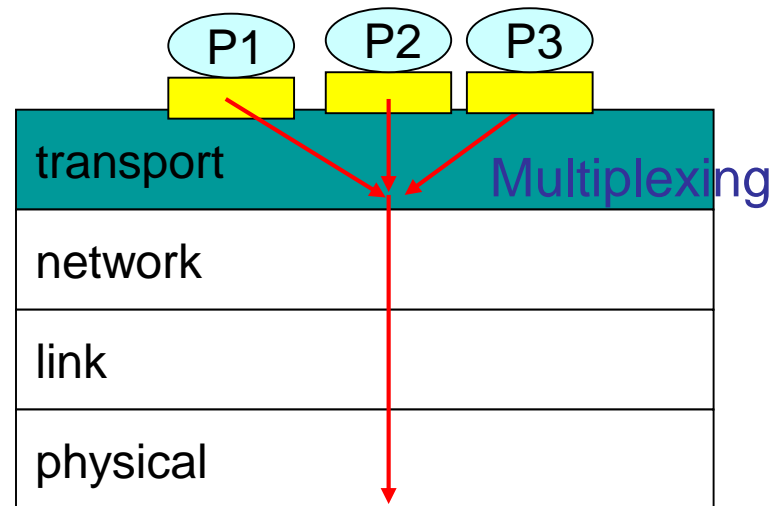
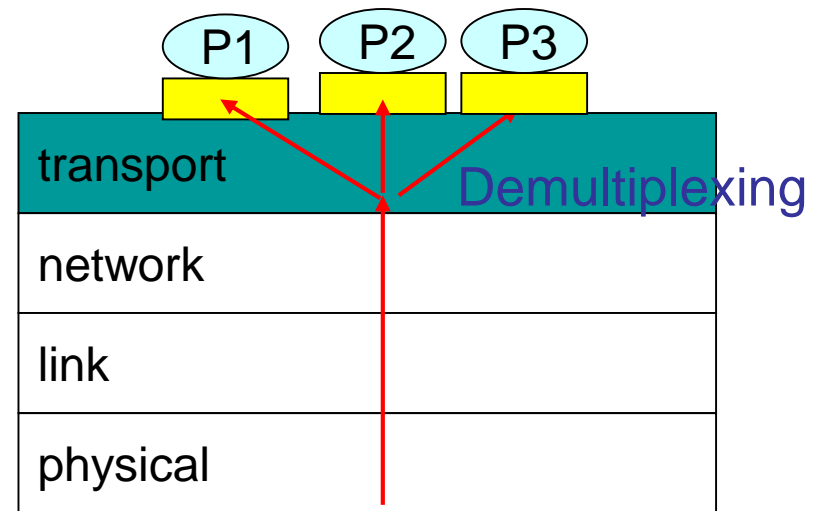
- 12 kids sending letters to 12 kids*
- app messages = letters in envelopes
  - processes = kids
  - hosts = houses
  - transport protocol = Ann and Bill
  - network-layer protocol = postal service

# Outline

- Transport-layer services
- Multiplexing and demultiplexing
- Connectionless transport: UDP

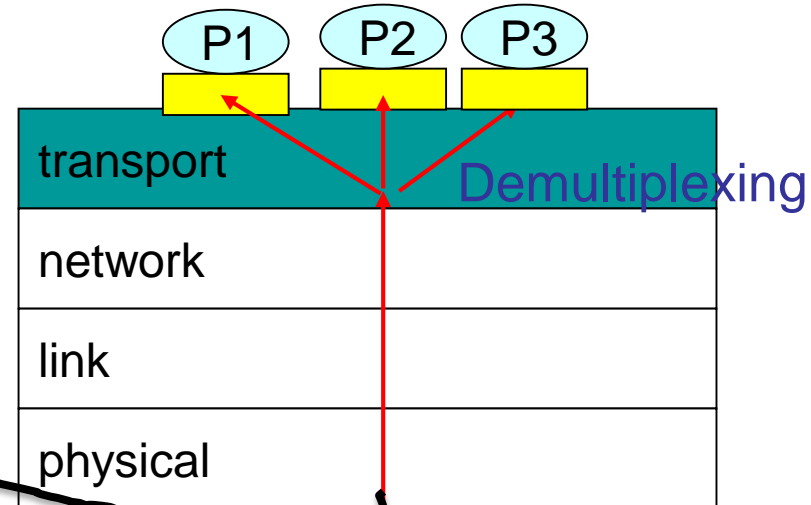
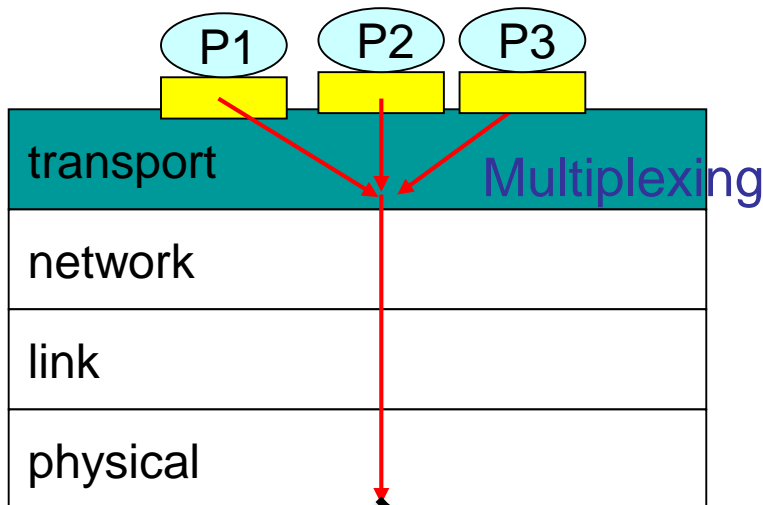
# Multiplexing/demultiplexing – Definition

- Demultiplexing
  - To deliver the data in a transport-layer segment to the correct socket
- Multiplexing
  - To create segments and pass them to the network layer.



# Multiplexing/demultiplexing – Why?

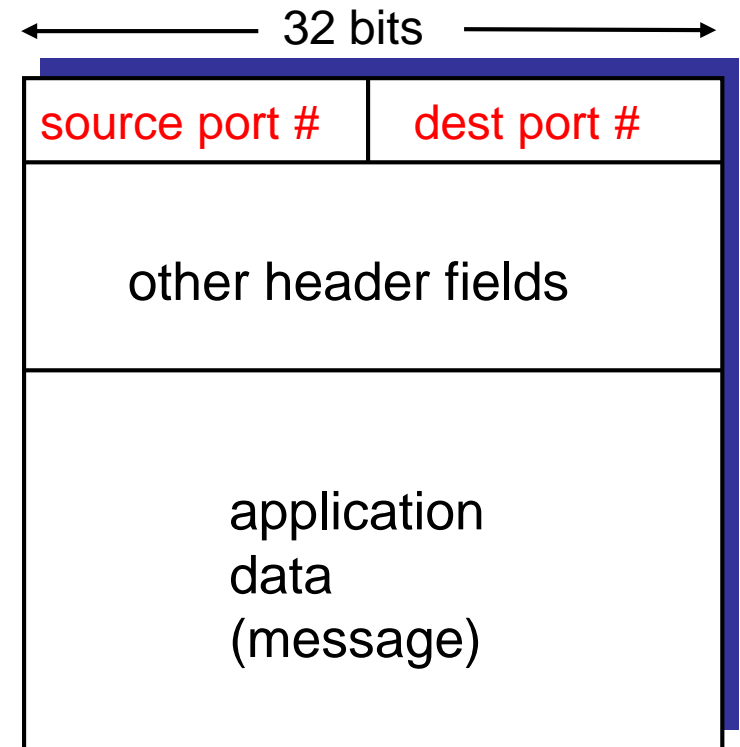
- Why need demultiplexing?
  - Now, the transport layer needs a way to determine which application the packet needs to be delivered. This is the **demultiplexing** problem.
- Why need multiplexing?
  - To facilitate demultiplexing!





# How demultiplexing works

- **host receives IP datagrams**
  - each datagram has source IP address, destination IP address
  - each datagram carries 1 transport-layer segment
  - each segment has source, destination port number
- **host uses IP addresses & port numbers to direct segment to appropriate socket**



TCP/UDP segment format

# Connectionless demultiplexing

- Create sockets with port numbers:

```
DatagramSocket mySocket1 = new  
    DatagramSocket(99111);
```

```
DatagramSocket mySocket2 = new  
    DatagramSocket(99222);
```

- UDP socket identified by two-tuple:

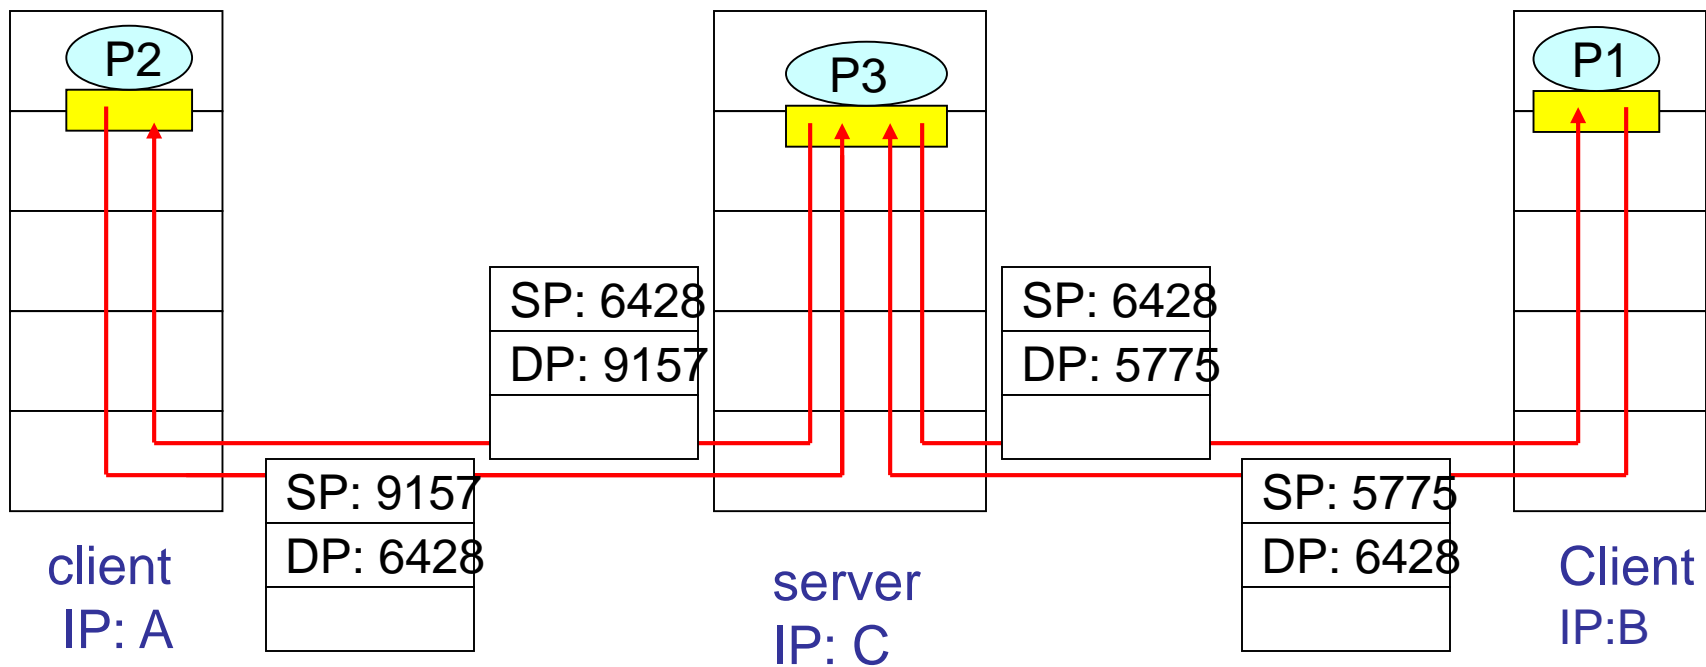
(dest IP address, dest port number)

Why do we need source port # ?

- When host receives UDP segment:
  - checks destination port number in segment
  - directs UDP segment to socket with that port number
- IP datagrams with different source IP addresses and/or source port numbers, but with the same destination IP and port number, will be directed to same socket.

# Connectionless demux (cont)

```
DatagramSocket serverSocket = new DatagramSocket(6428);
```

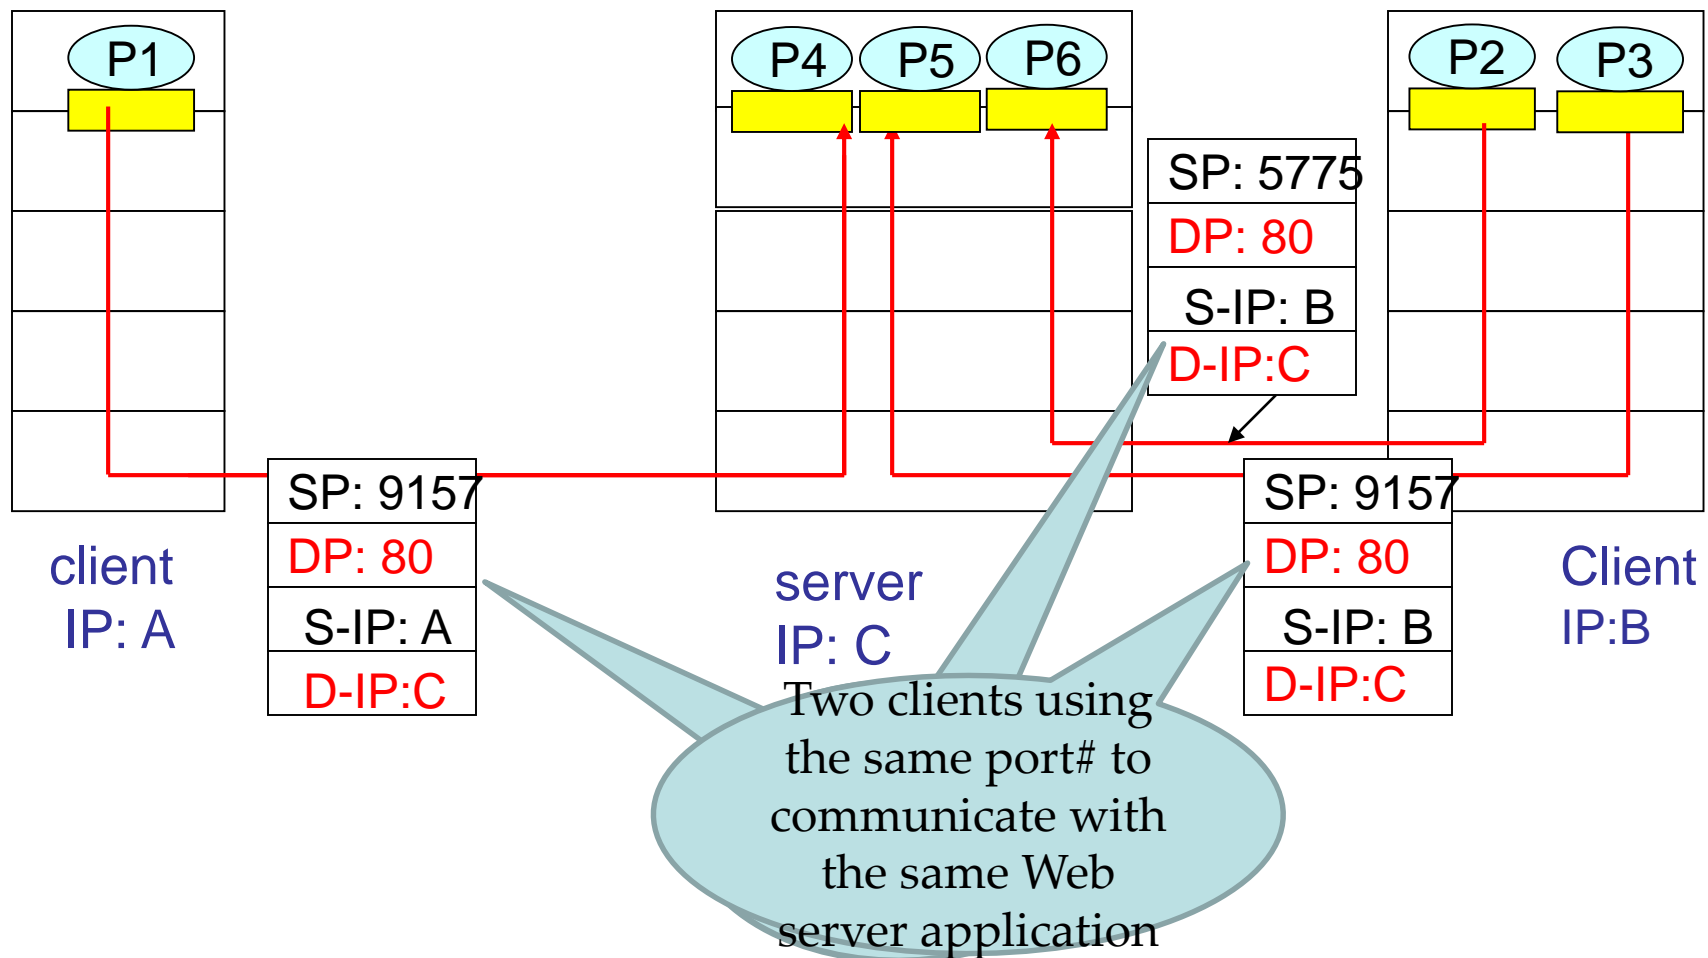


SP provides "return address"

# Connection-oriented demux

- TCP socket identified by 4-tuple:
  - source IP address
  - source port number
  - dest IP address
  - dest port number
- recv host uses all four values to direct segment to appropriate socket
- Server host may support many simultaneous TCP sockets:
  - each socket identified by its own 4-tuple
- Web servers have different sockets for each connecting client
  - non-persistent HTTP will have different socket for each request

# Connection-oriented demux (cont)



# Outline

- Transport-layer services
- Multiplexing and demultiplexing
- Connectionless transport: UDP

# UDP: User Datagram Protocol [RFC 768]

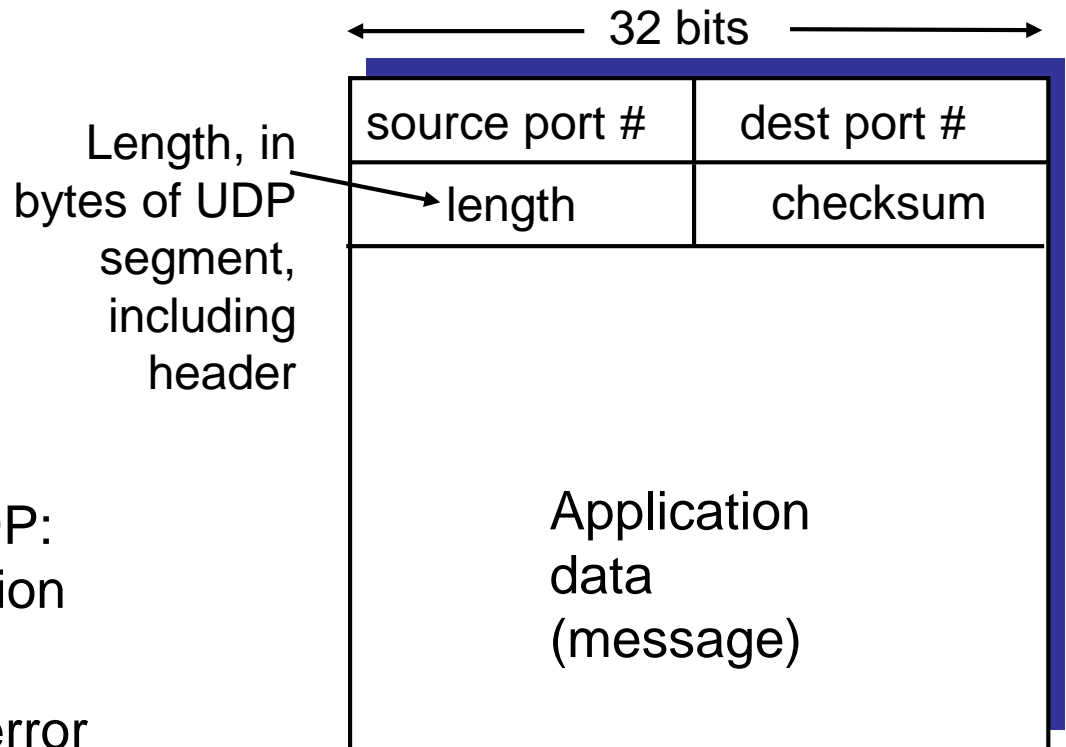
- “no frills,” “bare bones”  
Internet transport protocol
- “best effort” service, UDP segments may be:
  - lost
  - delivered out of order to app
- *connectionless*:
  - no handshaking between UDP sender, receiver
  - each UDP segment handled independently of others

## Why is there a UDP?

- no connection establishment (which can add delay)
- simple: no connection state at sender, receiver
- small segment header
- no congestion control: UDP can blast away as fast as desired

# UDP: more

- often used for streaming multimedia apps
  - loss tolerant
  - rate sensitive
- other UDP uses
  - DNS
  - SNMP
- reliable transfer over UDP: add reliability at application layer
  - application-specific error recovery!



UDP segment format



# Popular Internet Applications and their underlying transport protocols

Application	Application-layer protocol	Underlying transport protocol
Email	SMTP	TCP
Web	HTTP	TCP
Routing protocol	BGP	TCP
Routing protocol	RIP	Typically UDP
Name translation	DNS	Typically UDP
Streaming multimedia	Typically proprietary	Typically UDP
Internet telephony	Typically proprietary	Typically UDP

# UDP checksum

Goal: detect “errors” (e.g., flipped bits) in transmitted segment

## Sender:

- treat segment contents as sequence of 16-bit integers
- checksum: addition (1’s complement sum) of segment contents
- sender puts checksum value into UDP checksum field

## Receiver:

- compute checksum of received segment
- check if computed checksum equals checksum field value:
  - NO - error detected
  - YES - no error detected.  
*But maybe errors nonetheless? More later ....*

# Internet Checksum Example

- Note
  - When adding numbers, a carryout from the most significant bit needs to be added to the result
- Example: add two 16-bit integers

	1	1	1	0	0	1	1	0	0	1	1	0	0	1	1	0
	1	1	0	1	0	1	0	1	0	1	0	1	0	1	0	1
<hr/>																
wraparound	1	1	0	1	1	1	0	1	1	1	0	1	1	1	0	1
<hr/>																
sum	1	0	1	1	1	0	1	1	1	0	1	1	1	1	0	0
checksum	0	1	0	0	0	1	0	0	0	1	0	0	0	0	1	1

# A UDP segment captured by Wireshark

12	0.508950	10.34.40.169	132.181.2.225	DNS
13	0.509041	132.181.2.225	10.34.40.169	DNS
14	0.509462	10.34.40.169	132.181.2.225	DNS
15	0.510280	132.181.2.225	10.34.40.169	DNS
16	0.510354	132.181.2.225	10.34.40.169	DNS
17	33.503303	10.34.40.169	132.181.2.225	DNS
18	33.504666	132.181.2.225	10.34.40.169	DNS
19	61.002465	10.34.40.169	132.181.2.225	DNS
20	61.004045	132.181.2.225	10.34.40.169	DNS
21	61.013117	10.34.40.169	132.181.2.225	DNS
22	61.014578	132.181.2.225	10.34.40.169	DNS

- ▷ Frame 12: 88 bytes on wire (704 bits), 88 bytes captured (704 bits) on interface 0
- ▷ Ethernet II, Src: IntelCor\_b6:fe:63 (80:19:34:b6:fe:63), Dst: JuniperN\_ef:61:00 (2c:21:31:ef:61:00)
- ▷ Internet Protocol Version 4, Src: 10.34.40.169, Dst: 132.181.2.225
- ◀ User Datagram Protocol, Src Port: 63507, Dst Port: 53
  - Source Port: 63507
  - Destination Port: 53
  - Length: 54
  - Checksum: 0xe576 [unverified]
  - [Checksum Status: Unverified]
  - [Stream index: 6]
- ▷ Domain Name System (query)

# Why Would Anyone Use UDP?

- **Finer control over what data is sent and when**
  - As soon as an application process writes into the socket
  - ... UDP will package the data and send the packet
- **No delay for connection establishment**
  - UDP just blasts away without any formal preliminaries
  - ... which avoids introducing any unnecessary delays
- **No connection state**
  - No allocation of buffers, parameters, sequence #s, etc.
  - ... making it easier to handle many active clients at once
- **Small packet header overhead**
  - UDP header is only eight-bytes long

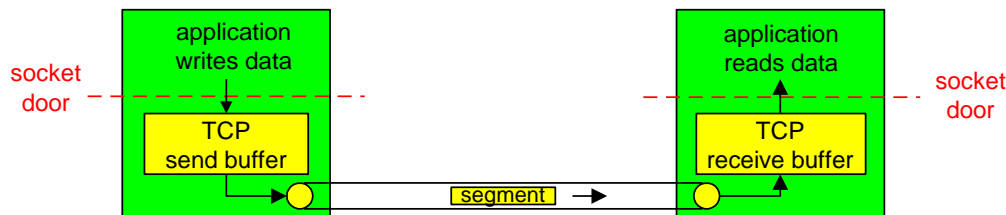
# Outline

- Connection-oriented transport: TCP
  - Overview and segment structure
  - Connection management
  - Reliable data transfer
  - Flow control

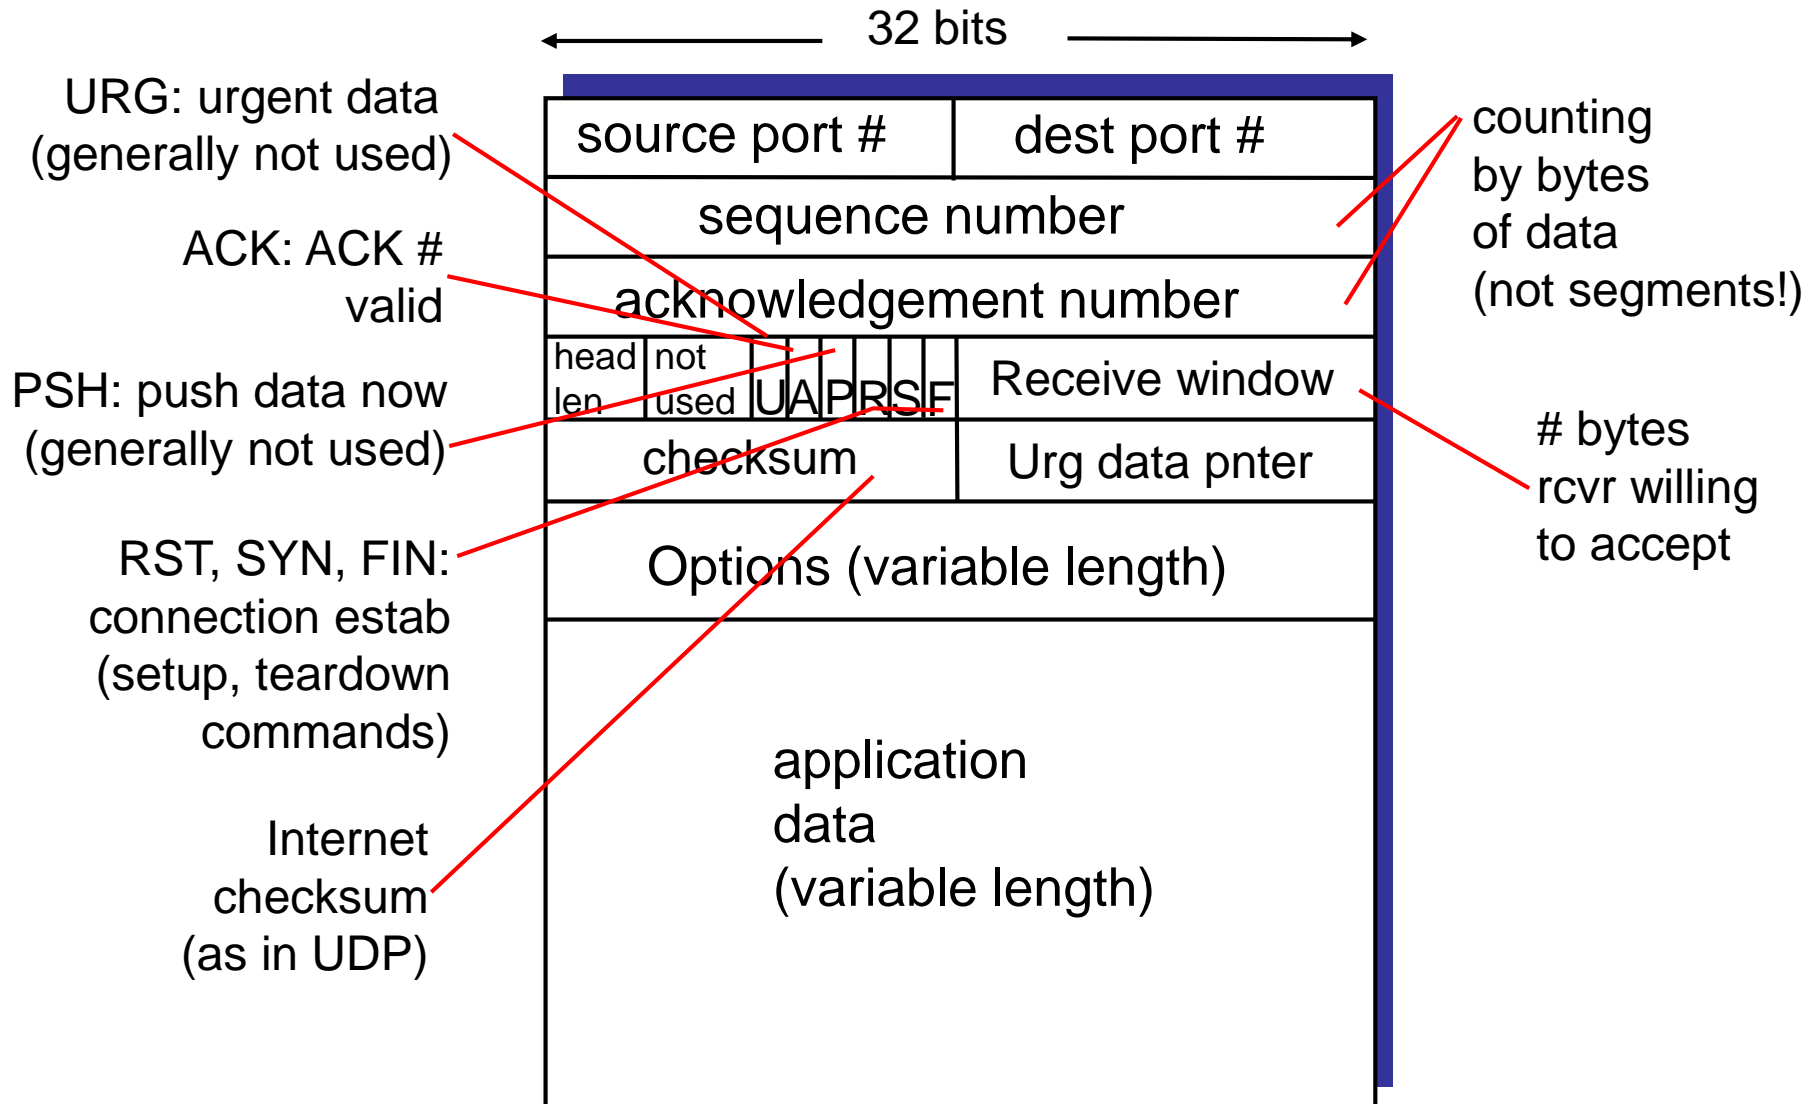
# TCP: Overview

RFCs: 793, 1122, 1323, 2018, 2581

- **point-to-point:**
  - one sender, one receiver
- **reliable, in-order *byte stream*:**
  - no “message boundaries”
- **pipelined:**
  - TCP congestion and flow control set window size
- ***send & receive buffers***
- **full duplex data:**
  - bi-directional data flow in same connection
  - MSS: maximum segment size
- **connection-oriented:**
  - handshaking (exchange of control msgs) init's sender, receiver state before data exchange
- **flow controlled:**
  - sender will not overwhelm receiver

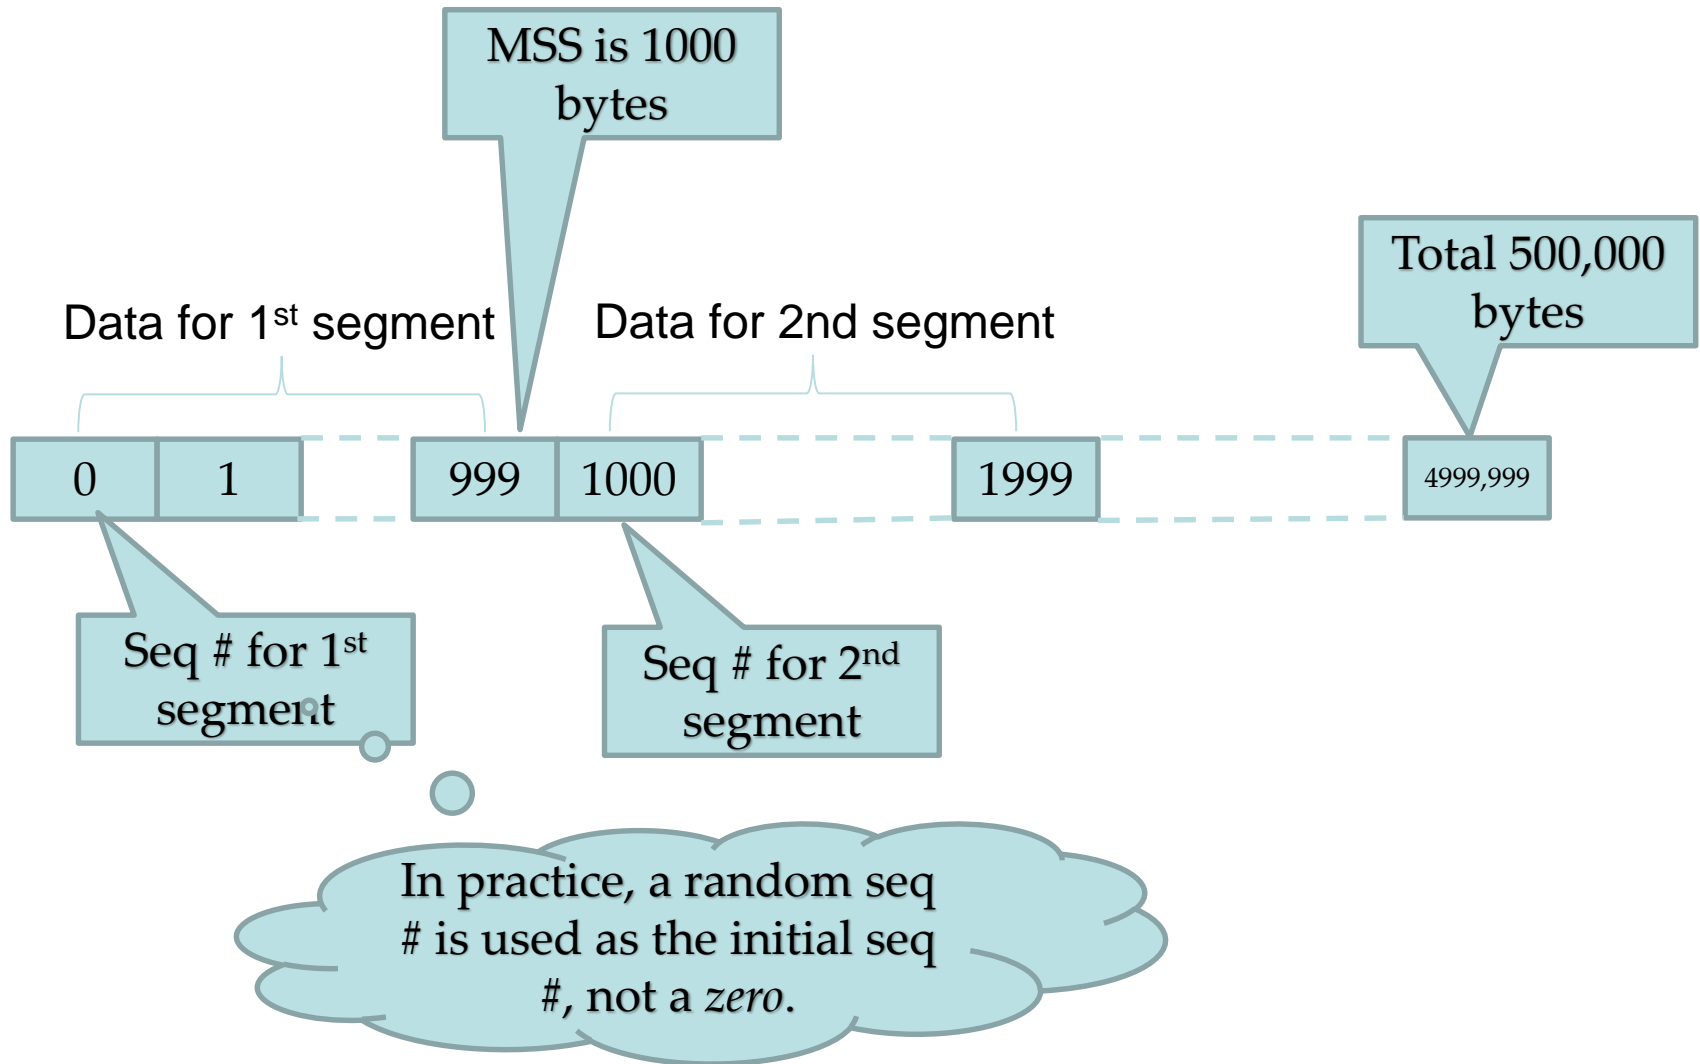


# TCP segment structure





# TCP segments



# TCP seq. #'s and ACKs

## Seq. #'s:

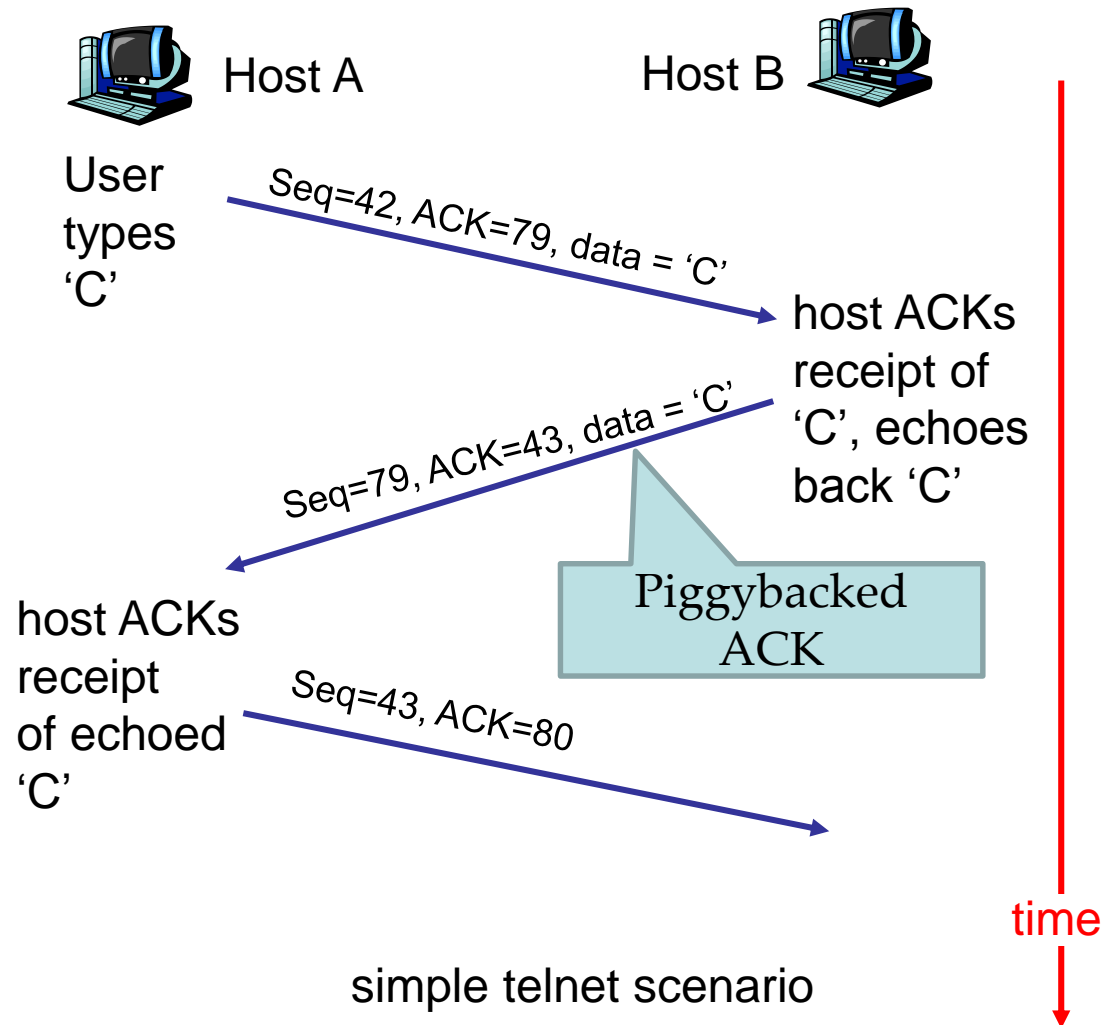
- byte stream  
“number” of first  
byte in segment's  
data

## ACKs:

- seq # of next byte  
expected from other  
side
- cumulative ACK

**Q:** how receiver handles  
out-of-order segments

- A: TCP spec doesn't  
say, - up to  
implementor



# TCP Connection Management

## Three way handshake:

Step 1: client host sends TCP SYN segment to server

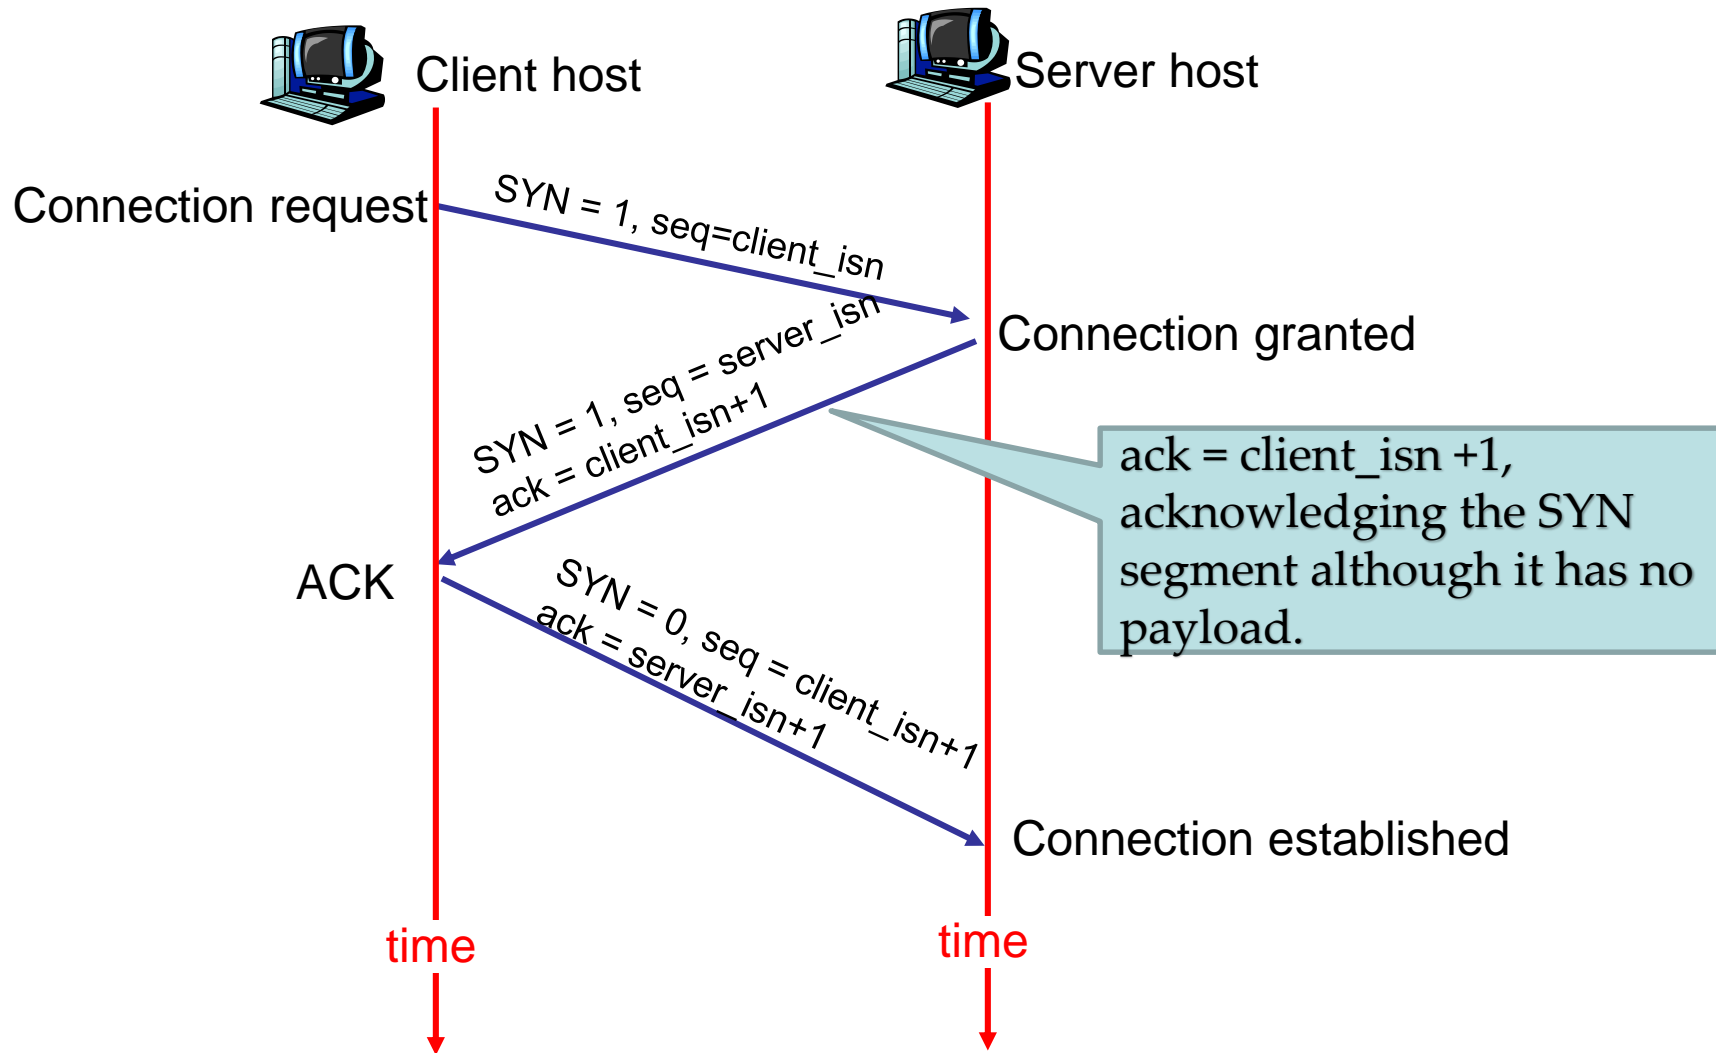
- specifies initial seq #
- no data

Step 2: server host receives SYN, replies with SYNACK segment

- server allocates buffers
- specifies server initial seq. #

Step 3: client receives SYNACK, replies with ACK segment, which may contain data

# TCP 3-way handshake



No.	Time	Source	Destination	Protocol
1	0.000000	192.168.88.161	128.238.26.21	TCP
2	0.217170	128.238.26.21	192.168.88.161	TCP
3	0.217326	192.168.88.161	128.238.26.21	TCP
4	51.906622	128.238.26.21	192.168.88.161	TCP
5	51.907091	192.168.88.161	128.238.26.21	TCP
6	51.907277	128.238.26.21	192.168.88.161	TCP

Frame 1: 74 bytes on wire (592 bits), 74 bytes captured (592 bits) on interface vmnic0

Ethernet II, Src: Vmware\_b9:06:32 (00:0c:29:b9:06:32), Dst: Vmware\_e5:ea:29 (00:0c:29:b9:06:29)

Internet Protocol Version 4, Src: 192.168.88.161, Dst: 128.238.26.21

Transmission Control Protocol, Src Port: 44058, Dst Port: 80

Source Port: 44058

Destination Port: 80

[Stream index: 0]

[TCP Segment Len: 0]

Sequence number: 2649682993

Acknowledgment number: 0

Header Length: 40 bytes

Flags: 0x002 (SYN)

Window size value: 29200

[Calculated window size: 29200]

Checksum: 0xcb45 [unverified]

[Checksum Status: Unverified]

Urgent pointer: 0

Options: (20 bytes), Maximum segment size 65535

Frame 2: 58 bytes on wire (464 bits), 58 bytes captured (464 bits) on interface vmnic0

Ethernet II, Src: Vmware\_e5:ea:29 (00:0c:29:b9:06:29), Dst: Vmware\_b9:06:32 (00:0c:29:b9:06:32)

Internet Protocol Version 4, Src: 128.238.26.21, Dst: 192.168.88.161

Transmission Control Protocol, Src Port: 80, Dst Port: 44058

Source Port: 80

Destination Port: 44058

[Stream index: 0]

[TCP Segment Len: 0]

Sequence number: 2058359514

Acknowledgment number: 2649682994

Header Length: 24 bytes

Flags: 0x012 (SYN, ACK)

Window size value: 64240

[Calculated window size: 64240]

Checksum: 0x1cc2 [unverified]

[Checksum Status: Unverified]

Urgent pointer: 0

Options: (4 bytes), Maximum segment size 65535

[SEQ/ACK analysis]

Frame 3: 54 bytes on wire (432 bits), 54 bytes captured (432 bits) on interface vmnic0

Ethernet II, Src: Vmware\_b9:06:32 (00:0c:29:b9:06:32), Dst: Vmware\_e5:ea:29 (00:0c:29:b9:06:29)

Internet Protocol Version 4, Src: 192.168.88.161, Dst: 128.238.26.21

Transmission Control Protocol, Src Port: 44058, Dst Port: 80, Seq: 2649682994, Len: 0

Source Port: 44058

Destination Port: 80

[Stream index: 0]

[TCP Segment Len: 0]

Sequence number: 2649682994

Acknowledgment number: 2058359515

Header Length: 20 bytes

Flags: 0x010 (ACK)

Window size value: 29200

[Calculated window size: 29200]

[Window size scaling factor: -2 (no window scaling used)]

Checksum: 0xbd5f [unverified]

[Checksum Status: Unverified]

Urgent pointer: 0

[SEQ/ACK analysis]

# TCP Connection Management (closing)

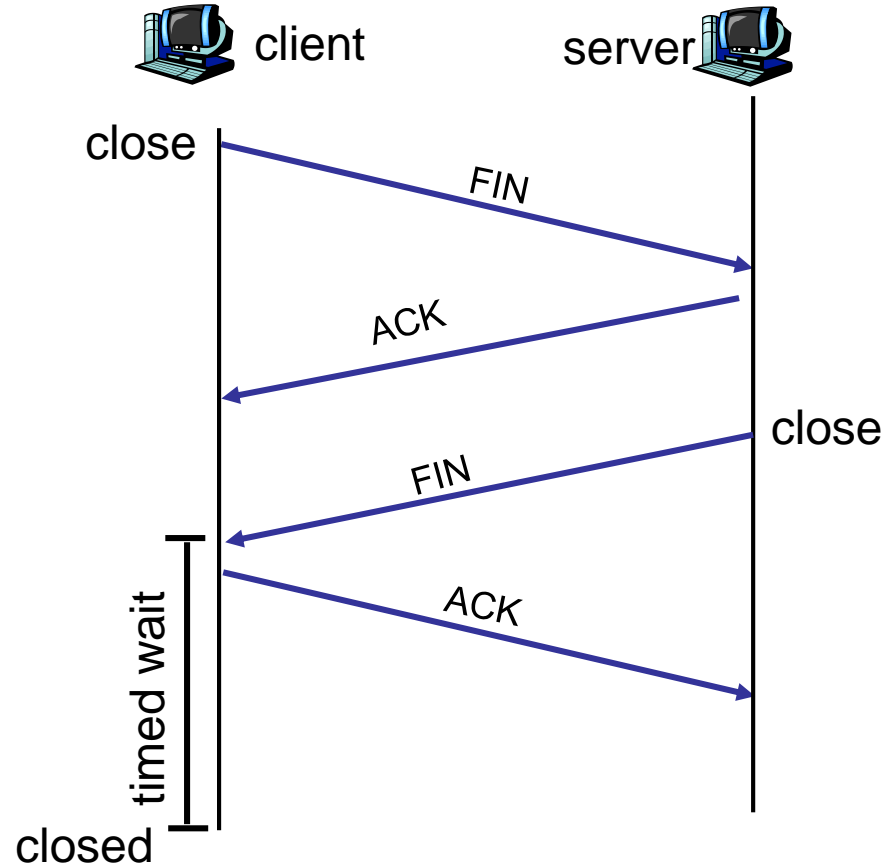
*All good things must come to an end:*

**Step 1:** client end system sends TCP FIN control segment to server

**Step 2:** server receives FIN, replies with ACK. Closes connection, sends FIN.

**Step 3:** client receives FIN, replies with ACK.

**Step 4:** server, receives ACK. Connection closed.



37	160.494102	10.34.40.169	132.181.107.25
38	160.495931	132.181.107.25	10.34.40.169
39	160.496329	132.181.107.25	10.34.40.169
40	160.496413	10.34.40.169	132.181.107.25

Internet Protocol Version 4, Src: 10.34.40.169, Dst: 132.181.107.25

Transmission Control Protocol, Src Port: 57483, Dst Port: 25, Seq: 3400705033, Ack: 1545042735,

Source Port: 57483  
Destination Port: 25  
[Stream index: 1]  
[TCP Segment Len: 0]  
Sequence number: 3400705033  
Acknowledgment number: 1545042735  
Header Length: 20 bytes

Flags: 0x011 (FIN, ACK)

000. .... = Reserved: Not set  
...0 .... = Nonce: Not set  
...0... .... = Congestion Window Reduced  
...0... .... = ECN-Echo: Not set  
...0... .... = Urgent: Not set  
...0...1 .... = Acknowledgment: Set  
...0...0... = Push: Not set  
...0...0... = Reset: Not set  
...0...0... = Syn: Not set

...1 = Fin: Set

[TCP Flags: .....A...F]

Window size value: 256

[Calculated window size: 65536]

[Window size scaling factor: 256]

Checksum: 0x6592 [unverified]

[Checksum Status: Unverified]

Urgent pointer: 0

FIN

Frame 38: 54 bytes on wire (432 bits)

Ethernet II, Src: JuniperN\_ef:61:00 (

Internet Protocol Version 4, Src: 132

Transmission Control Protocol, Src Po

Source Port: 25

Destination Port: 57483

[Stream index: 1]

[TCP Segment Len: 0]

Sequence number: 1545042735

Acknowledgment number: 3400705034

Header Length: 20 bytes

Flags: 0x010 (ACK)

000. .... = Reserved: Not se

...0 .... = Nonce: Not set

...0... .... = Congestion Windc

...0... .... = ECN-Echo: Not set

...0... .... = Urgent: Not set

...0...1 .... = Acknowledgment: Set

...0...0... = Push: Not set

...0...0... = Reset: Not set

...0...0... = Syn: Not set

...0...0... = Fin: Not set

[TCP Flags: .....A....]

Window size value: 513

[Calculated window size: 131328]

[Window size scaling factor: 256]

Checksum: 0x6491 [unverified]

ACK

37	160.494102	10.34.40.169	132.181.107.25
38	160.495931	132.181.107.25	10.34.40.169
39	160.496329	132.181.107.25	10.34.40.169
40	160.496413	10.34.40.169	132.181.107.25

Frame 40: 54 bytes on wire (432 bits), 54 bytes captured (432 bits) on interface 0

Ethernet II, Src: IntelCor\_b6:fe:63 (80:19:34:b6:fe:63), Dst: JuniperN\_ef:61:00 (2c:21:31:ef

Internet Protocol Version 4, Src: 10.34.40.169, Dst: 132.181.107.25

Transmission Control Protocol, Src Port: 57483, Dst Port: 25, Seq: 3400705034, Ack: 15450427

Source Port: 57483

Destination Port: 25

[Stream index: 1]

[TCP Segment Len: 0]

Sequence number: 3400705034

Acknowledgment number: 1545042736

Header Length: 20 bytes

Flags: 0x010 (ACK)

000. .... = Reserved: Not set

...0 .... = Nonce: Not set

...0... .... = Congestion Window Reduced (CWR): Not set

...0... .... = ECN-Echo: Not set

...0... .... = Urgent: Not set

...0...1 .... = Acknowledgment: Set

...0...0... = Push: Not set

...0...0... = Reset: Not set

...0...0... = Syn: Not set

...0...0... = Fin: Not set

[TCP Flags: .....A....]

Window size value: 256

[Calculated window size: 65536]

[Window size scaling factor: 256]

Checksum: 0x6591 [unverified]

[TCP Segment Len: 0]

Sequence number: 1545042735

Acknowledgment number: 3400705034

Header Length: 20 bytes

Flags: 0x011 (FIN, ACK)

000. .... = Reserved: Not set

...0 .... = Nonce: Not set

...0... .... = Congestion Window Reduced (CWR): Not set

...0... .... = ECN-Echo: Not set

...0... .... = Urgent: Not set

...0...1 .... = Acknowledgment: Set

...0...0... = Push: Not set

...0...0... = Reset: Not set

...0...0... = Syn: Not set

...0...1 = Fin: Set

[TCP Flags: .....A...F]

Window size value: 513

[Calculated window size: 131328]

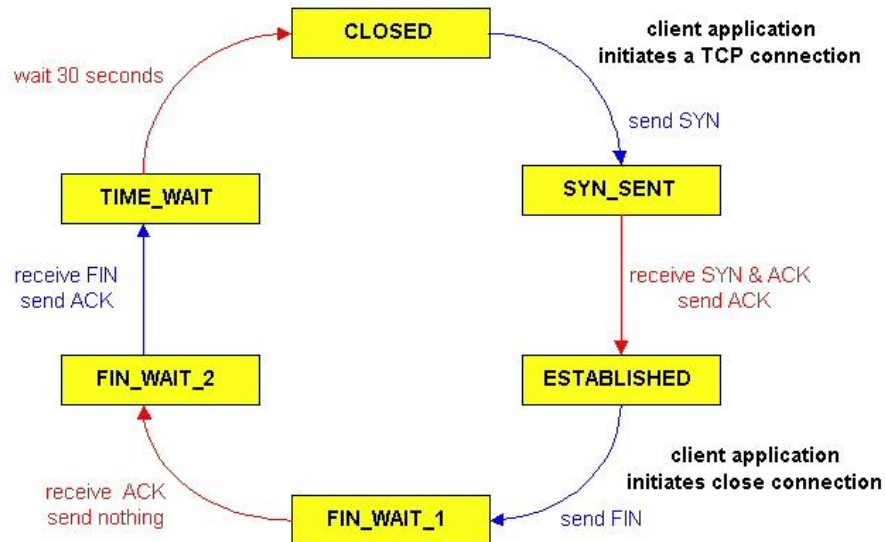
[Window size scaling factor: 256]

Checksum: 0x6490 [unverified]

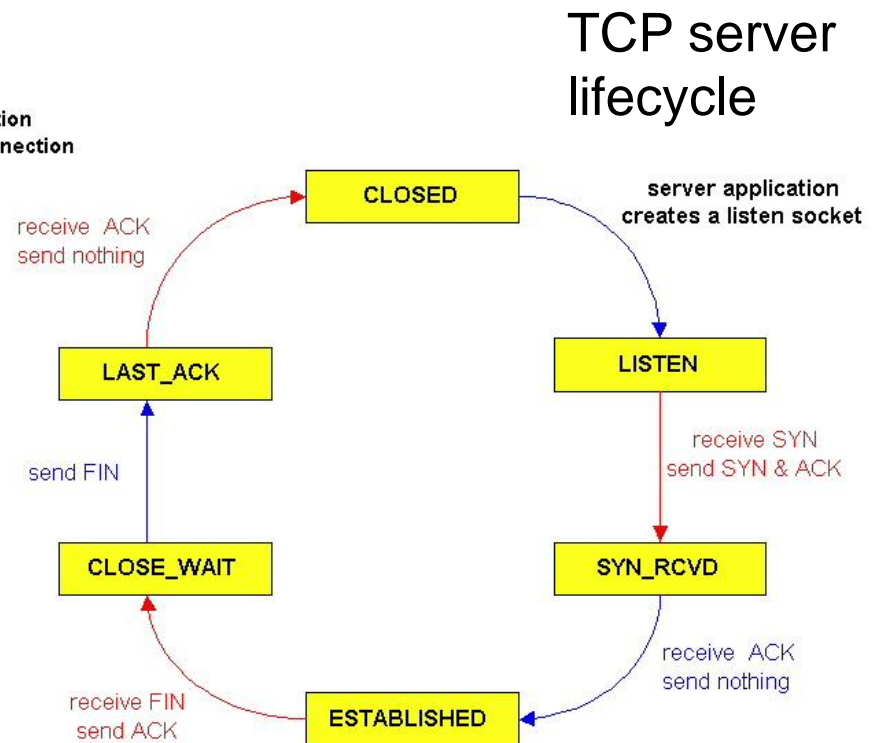
FIN

ACK

# TCP Connection Management (cont)



TCP client lifecycle



TCP server lifecycle



# Outline

- Connection-oriented transport: TCP
  - Overview and segment structure
  - Connection management
  - **Reliable data transfer**
  - Flow control

# References

- [KR3] James F. Kurose, Keith W. Ross, *Computer networking: a top-down approach featuring the Internet*, 3<sup>rd</sup> edition.
- [PD5] Larry L. Peterson, Bruce S. Davie, *Computer networks: a systems approach*, 5<sup>th</sup> edition
- [TW5] Andrew S. Tanenbaum, David J. Wetherall, *Computer network*, 5<sup>th</sup> edition
- [LHBi]Y-D. Lin, R-H. Hwang, F. Baker, *Computer network: an open source approach*, International edition

# Acknowledgements

- All slides are developed based on slides from the following two sources:
  - Dr DongSeong Kim's slides for COSC264, University of Canterbury;
  - Prof Aleksandar Kuzmanovic's lecture notes for CS340, Northwestern University,  
[https://users.cs.northwestern.edu/~akuzma/classes/CS340-w05/lecture\\_notes.htm](https://users.cs.northwestern.edu/~akuzma/classes/CS340-w05/lecture_notes.htm)