

COSC264

Introduction to Computer Networks and the Internet

# Introduction to the Web and HTTP

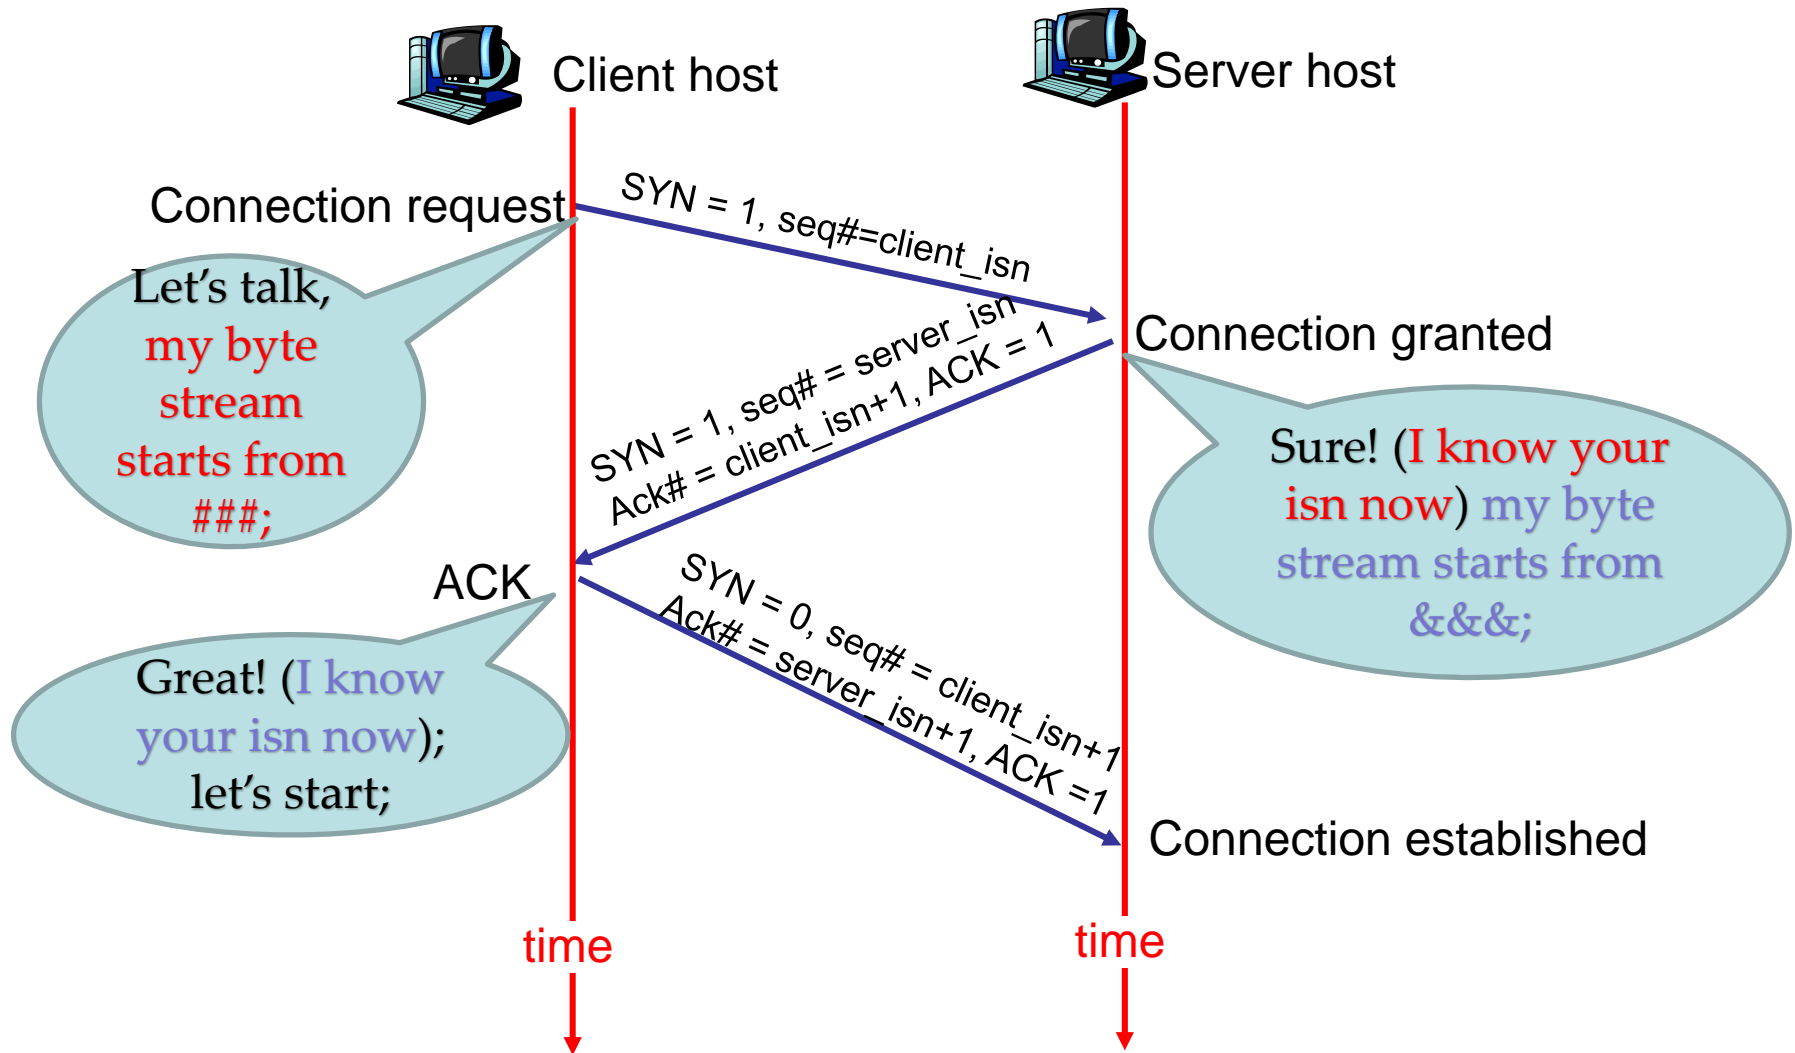
Dr. Barry Wu

Wireless Research Centre

University of Canterbury

[barry.wu@canterbury.ac.nz](mailto:barry.wu@canterbury.ac.nz)

# Why 3-way handshake? Not 2-way handshake?



# Outline

- Network applications
  - Network apps vs app. protocols
  - Application structure
- The Web
- HTTP

# Outline

- Network applications
  - Network apps vs app. protocols
  - Application structure
- The Web
- HTTP

# The most popular network applications

- The Web
  - Search engine
  - E-commerce
  - Banking
  - News
  - Online video
  - Blog
- Email
- Instant messaging
  - Text/pic/sound
  - Video call
- Video conferencing/gaming

# Other applications

- telnet
- File transfer (ftp)
- News group
- P2P file sharing

# What is a network application?

- Programs that run on different end system and communicate with each other over the network;
  - Web browser program in the user's host;
  - Web server program in the web server host;

# Network apps vs application protocols


- An application protocol is only one piece (a big one) of a network application.

Network app	Application protocol
Web (other pieces: HTML, web browsers, web servers)	HTTP (one piece of web, protocol)
Email	SMTP



# Services an app needs

- Reliable data transfer
  - Email; instant messaging; file transfer, financial applications;
  - There are loss-tolerant applications though;
    - Multimedia applications



# Bandwidth

- Rate
  - Internet telephony application
  - Many multimedia apps are bandwidth-sensitive; (*adaptive coding technique*)
- Elastic apps
  - Not strict with bandwidth;
  - Email; file transfer;

# Timing

- Tight timing constraints
  - Interactive real-time apps: Internet telephony, virtual environments (VR), teleconferencing, multiplayer games;
  - End-to-end delay:  $< 100\text{s of ms}$ ;

# Services provided by the Internet transport layer

TCP	UDP
Connected-oriented service	Connectionless service
Reliable transport service	<i>Unreliable data transfer service</i>
Congestion control ( <i>NO guarantee of min transmission rate</i> )	<i>No congestion control</i>
Flow control	<i>No flow control</i>
<i>No guarantee of delay</i>	<i>No guarantee of timing</i>

The Internet has been hosting time-sensitive applications for many years!


# Outline

- Network applications
  - Network apps vs app. protocols
  - Application structure
- The Web
- HTTP

# Application architectures

- Client-server
- Peer-to-peer (P2P)
- Hybrid of client-server and P2P

# Client-server architecture



## server:

- always-on host
- permanent IP address
- server farms for scaling

## clients:


- communicate with server
- may be intermittently connected
- may have dynamic IP addresses
- do not communicate directly with each other

# Pure P2P architecture

- no always on server
- arbitrary end systems directly communicate
- peers are intermittently connected and change IP addresses
- example: Gnutella

Highly scalable

But difficult to manage





# Hybrid of client-server and P2P

## Napster

- File transfer P2P
- File search centralized:
  - Peers register content at central server
  - Peers query same central server to locate content

## Instant messaging

- Chatting between two users is P2P
- Presence detection/location centralized:
  - User registers its IP address with central server when it comes online
  - User contacts central server to find IP addresses of buddies

# Outline

- Network applications
  - Network apps vs app. protocols
  - Application structure
- The Web
- HTTP

# The Web

- Until 1990s the Internet was mainly used by researchers, academics and university students.

## NAME

**telnet** – user interface to the TELNET protocol

## SYNOPSIS

**telnet** [-468ELadr] [-S tos] [-b address] [-e escapechar] [-l user]  
[-n tracefile] [host [port]]

## DESCRIPTION

The **telnet** command provides interactive communication with another host using the TELNET protocol. It can be run in command mode, where it prints a telnet prompt, or in user mode, where it is invoked with a host argument, in which case it attempts to connect to the host specified below.

Not everyone can use this!

## NAME

**ftp** – Internet file transfer program

## SYNOPSIS

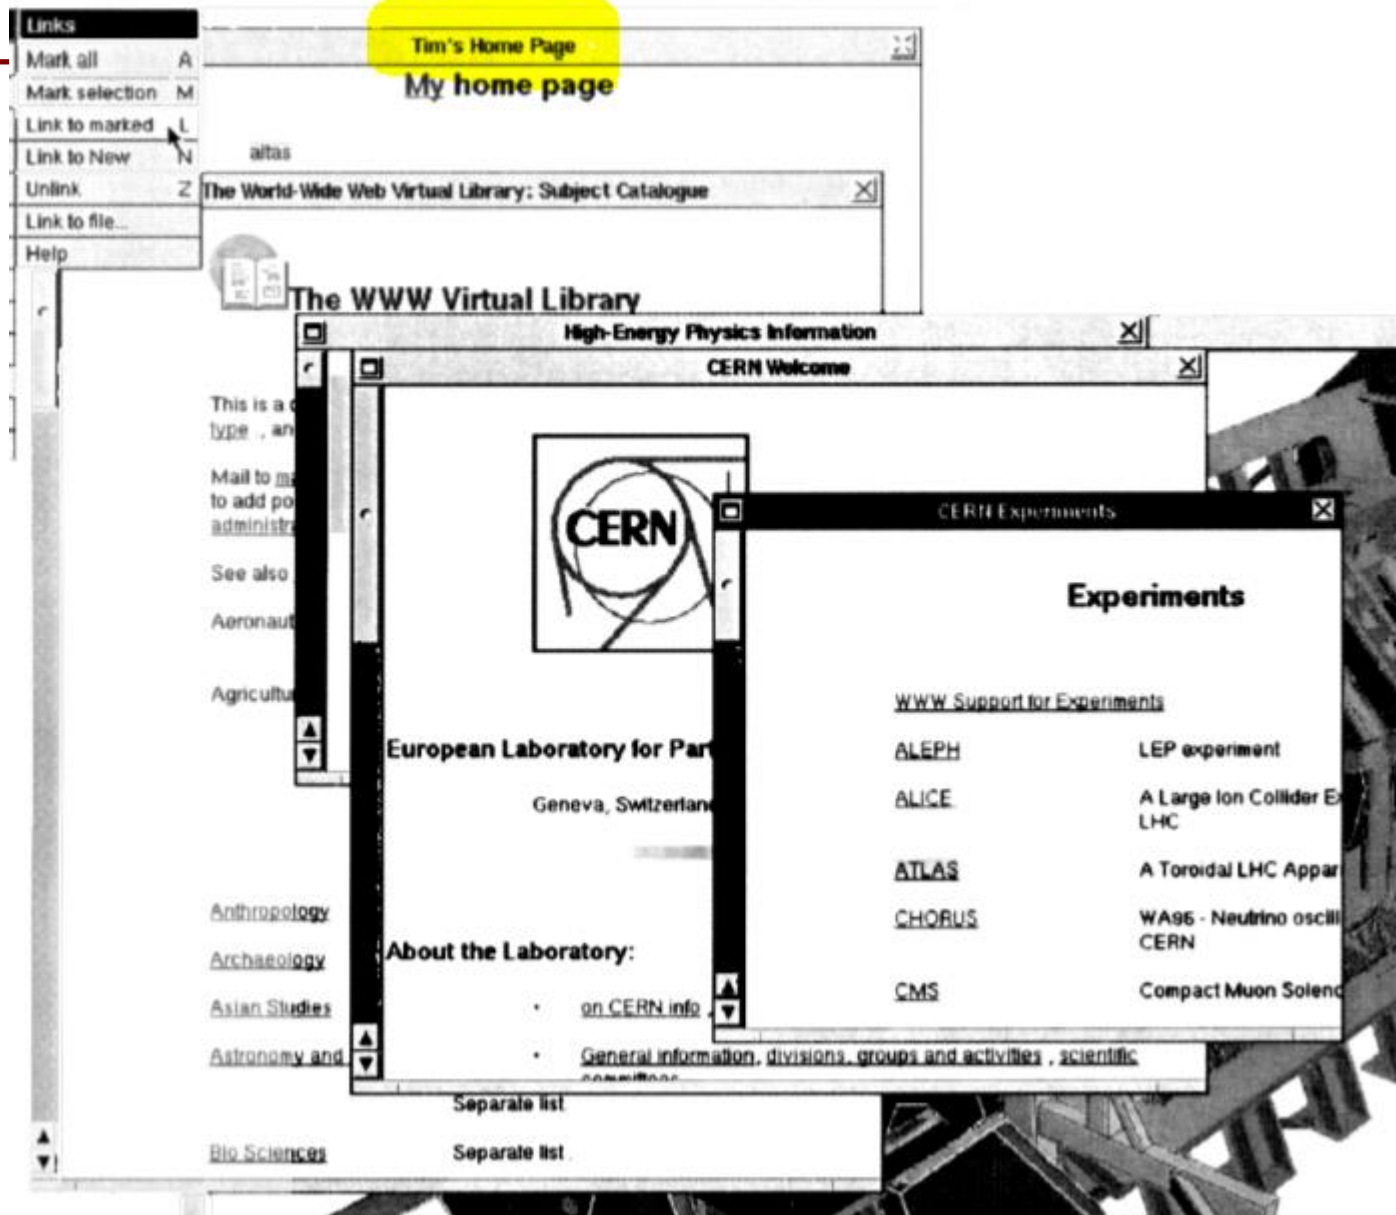
**ftp** [-46pinegvd] [host [port]]  
**pftp** [-46inegvd] [host [port]]

## DESCRIPTION

**Ftp** is the user interface to the Internet standard File Transfer Protocol. The program allows a user to transfer files to and from a remote network site.

Options may be specified at the command line, or to the command interpreter.

- Started by Tim Berners-Lee in 1989.
  - Telephone (1870s); radio (1920s)/TV (1930s);
  - Email and the Web;



- On-demand service!
- HTTP (HyperText Transfer Protocol) is at the heart of the Web.

# Outline

- Network applications
  - Network apps vs app. protocols
  - Application structure
- The Web
- HTTP



# Outline

- Network applications
  - Network apps vs app. protocols
  - Application structure
- The Web
- HTTP
  - Overview
  - Non-persistent HTTP and persistent HTTP
  - HTTP messages

# HTTP overview

## HTTP: hypertext transfer protocol

- Web's application layer protocol
- HTTP /1.0: RFC 1945
- HTTP /1.1: RFC 2068/2616/7230
- HTTP /2: RFC 7540
- HTTP /3: Introduced in *26 September 2019*.

# Web page

- **Web page** consists of **objects**
- Object can be HTML file, JPEG image, audio file,...
- Web page consists of **base HTML-file** which includes several referenced objects
- Each object is addressable by a **URL**
- Example URL:

Method

`http://www.someschool.edu/someDept/pic.gif`


host name

path name

## client/server model

*client:* browser that requests, receives, “displays” Web objects

*server:* Web server sends objects in response to requests



# HTTP overview (continued)

## Uses TCP:

- client initiates TCP connection (creates socket) to server, **port 80**
- server accepts TCP connection from client
- HTTP messages (application-layer protocol messages) exchanged between browser (HTTP client) and Web server (HTTP server)
- TCP connection closed

## HTTP is “stateless”

- server maintains no information about past client requests

aside  
Protocols that maintain “state” are complex!

- \* past history (state) must be maintained
- \* if server/client crashes, their views of “state” may be inconsistent, must be reconciled

# HTTP connections

## Nonpersistent HTTP

- At most *one object* is sent over a TCP connection.
- HTTP/1.0 uses nonpersistent HTTP

## Persistent HTTP

- Multiple objects can be sent over single TCP connection between client and server.
- HTTP/1.1 uses persistent connections in default mode

# Nonpersistent HTTP

Suppose user enters URL

`www.someSchool.edu/someDepartment/home.index`

(contains text,  
references to 10  
jpeg images)

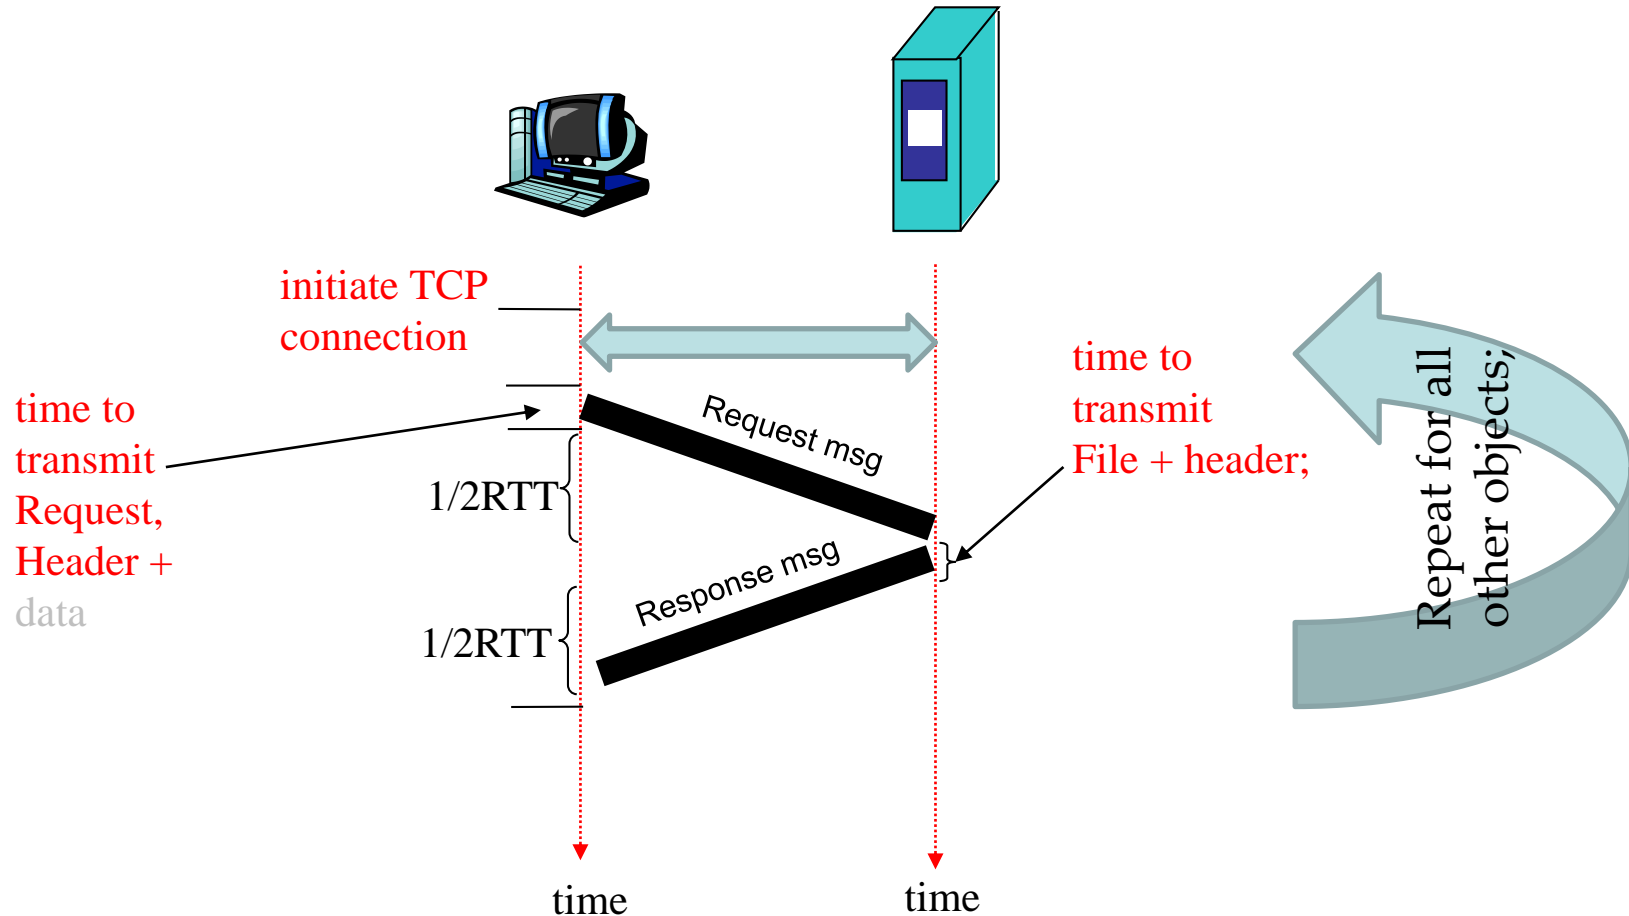
1. HTTP client initiates TCP connection to HTTP server (process) at `www.someSchool.edu` on port 80 (default port #)
2. HTTP client sends HTTP *request message* (containing URL) into TCP connection socket. Message indicates that client wants object `someDepartment/home.index`
3. HTTP server receives request message, forms *response message* containing requested object, and sends message into its socket

# Nonpersistent HTTP (cont.)

4. HTTP server process tells TCP to close the TCP connection.  
(It doesn't actually close it until it knows for sure the client has received the response message.)
5. HTTP client receives response message containing html file, displays html. Parsing html file, finds 10 referenced jpeg objects
6. Steps 1-5 repeated for each of 10 jpeg objects



# A simplified time modelling




# Another response time modeling

**Definition of RRT:** time to send a small packet to travel from client to server and back.

## Response time:

- *one RTT to initiate TCP connection*
- one RTT for HTTP request and first few bytes of HTTP response to return
- file transmission time

**total =  $2RTT + \text{transmit time}$**



Non-persistent HTTP requires reconnection for every objects.

# Persistent HTTP

## Nonpersistent HTTP issues:

- requires 2 RTTs per object
- OS must work and allocate host resources for each TCP connection
- *but browsers often open parallel TCP connections to fetch referenced objects*

## Persistent HTTP

- server leaves connection open after sending response
- subsequent HTTP messages between same client/server are sent over connection

# Persistent HTTP + pipelining

## Persistent without pipelining:

- client issues new request only when previous response has been received
- one RTT for each referenced object

## Persistent with pipelining:

- default in HTTP/1.1
- client sends requests as soon as it encounters a referenced object
- as little as one RTT for all the referenced objects

Nonpersistent HTTP    Using parallel TCP connections



Less response time

persistent HTTP without pipelining

persistent HTTP with pipelining

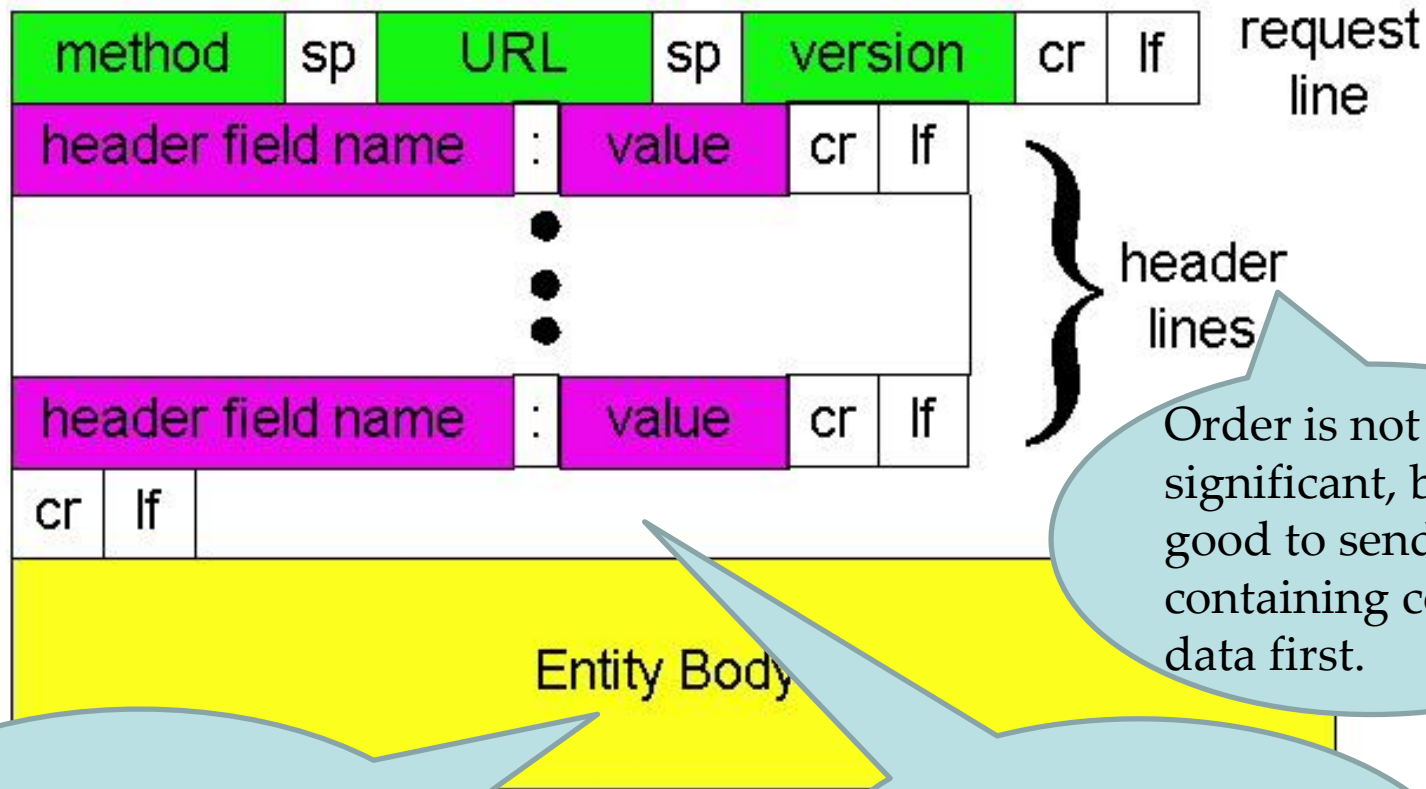
# Outline

- Network applications
  - Network apps vs app. protocols
  - Application structure
- The Web
- HTTP
  - Overview
  - Non-persistent HTTP and persistent HTTP
  - HTTP messages

# HTTP Message Format

two types of HTTP messages: *request, response*

# HTTP request message: general format



Order is not significant, but it is good to send fields containing control data first.

Empty for GET but it is used with POST (filling a form);

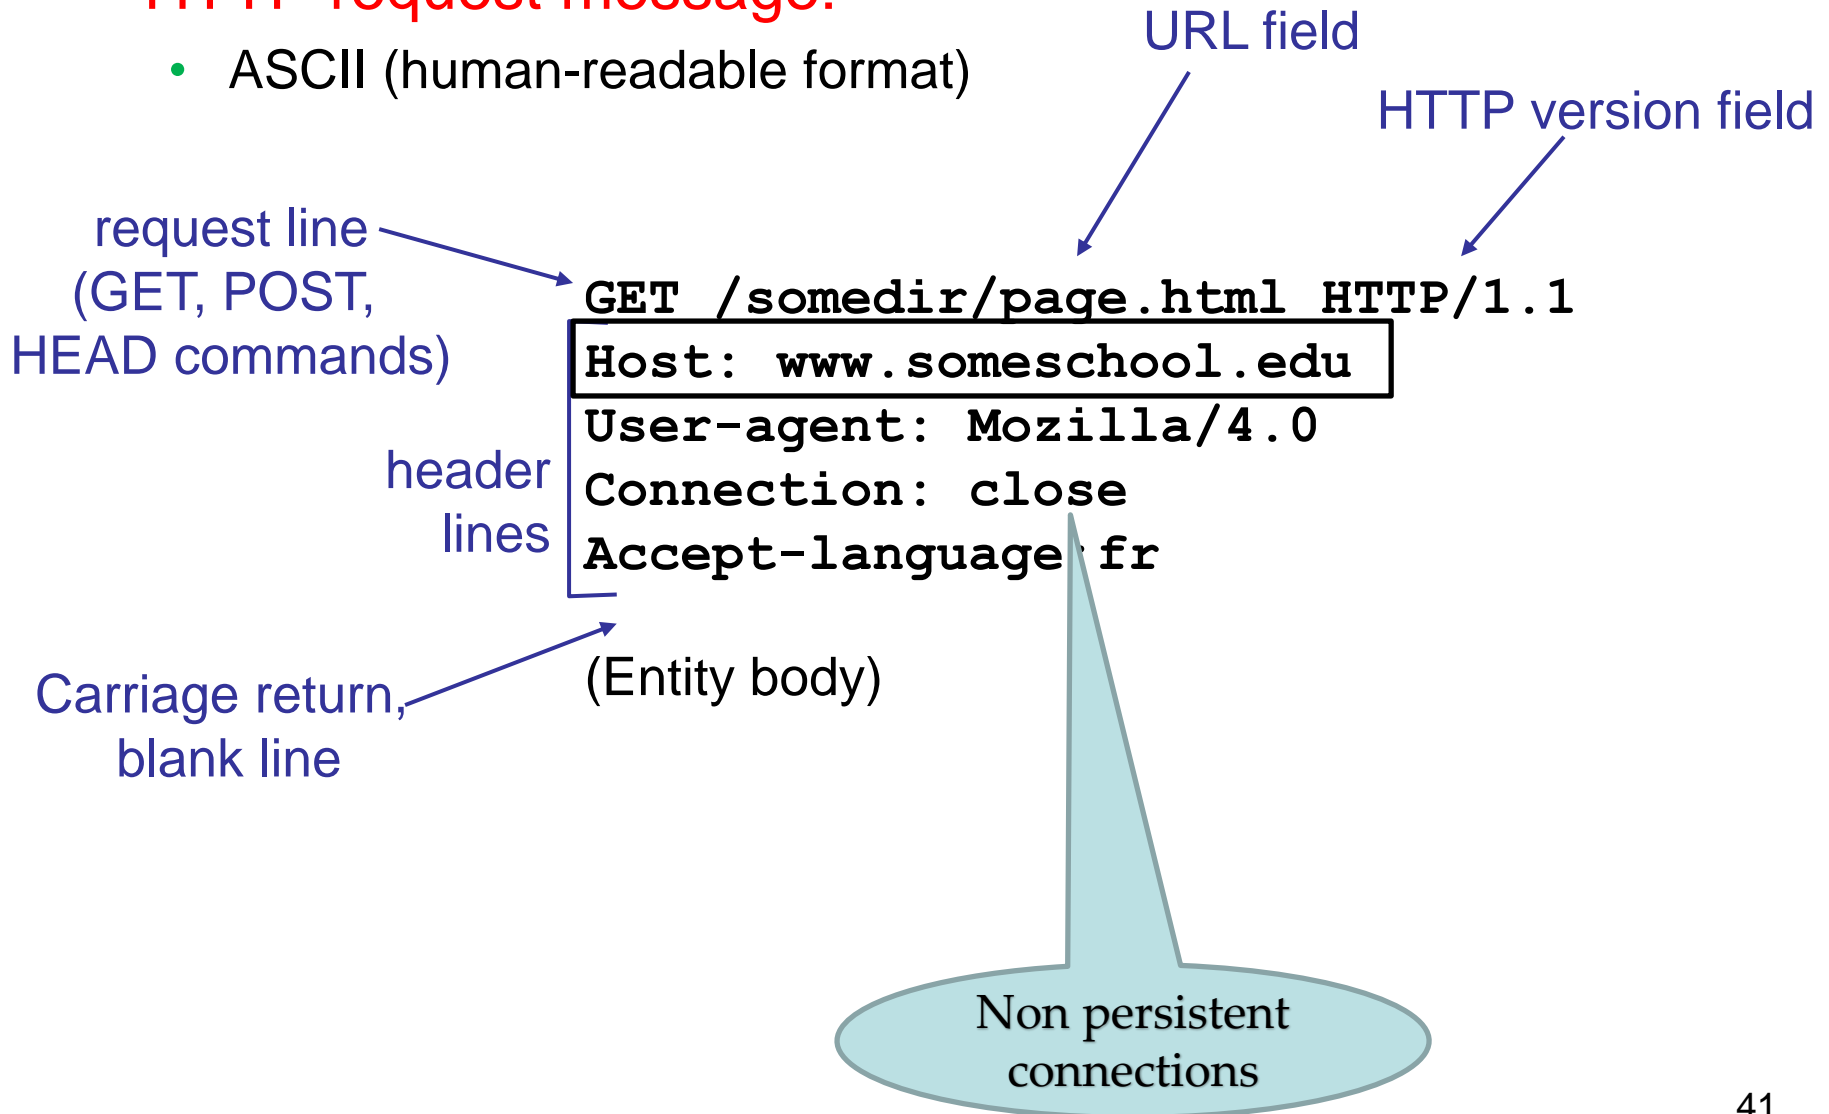
An empty line to indicate the end of the header section



# HTTP request message

- **HTTP request message:**

- ASCII (human-readable format)



Each language-range MAY be given an associated quality value which represents an estimate of the user's preference for the languages specified by that range. The quality value defaults to "q=1".

For example,


Accept-Language: da, en-gb;q=0.8, en;q=0.7

would mean: *"I prefer Danish, but will accept British English and other types of English."*

Form submission can be done with GET as well;  
“www.somesite.com/animalsearch?monkeys&bananas”

First name:

Last name:



## Submitted Form Data

Your input was received as:

```
fname=Zak&lname=Johnson
```

The server has processed your input and returned this answer.

# Method types


## HTTP/1.0

- GET
- POST – filling a form
- HEAD
  - asks server to leave requested object out of response
  - Can be used for debugging

## HTTP/1.1

- GET, POST, HEAD
- PUT
  - uploads file in entity body to path specified in URL field
- DELETE
  - deletes file specified in the URL field

# General format of a response message



# HTTP response status codes

## 200 OK

- request succeeded, requested object later in this message

## 301 Moved Permanently

- requested object moved, new location specified later in this message (Location:)

## 400 Bad Request

- request message not understood by server

## 404 Not Found

- requested document not found on this server

## 505 HTTP Version Not Supported

# HTTP response message

status line  
(protocol ver,  
status code,  
status msg.)

header  
lines

data, e.g.,  
requested  
HTML file

```
HTTP/1.1 200 OK
Connection close
Date: Thu, 06 Aug 1998 12:00:15 GMT
Server: Apache/1.3.0 (Unix)
Last-Modified: Mon, 22 Jun 1998 .....
Content-Length: 6821
Content-Type: text/html

data data data data data ...
```

The diagram illustrates the structure of an HTTP response message. It consists of three main parts: a status line, header lines, and data. The status line is labeled 'status line (protocol ver, status code, status msg.)' and points to the first line of the message, 'HTTP/1.1 200 OK'. The header lines are labeled 'header lines' and point to the subsequent lines: 'Connection close', 'Date: Thu, 06 Aug 1998 12:00:15 GMT', 'Server: Apache/1.3.0 (Unix)', 'Last-Modified: Mon, 22 Jun 1998 .....', 'Content-Length: 6821', and 'Content-Type: text/html'. The data part is labeled 'data, e.g., requested HTML file' and points to the final line, 'data data data data data ...'.

```
duser@192.168.88.155:~/libbgpdump-1.4.99.11$ telnet cis.poly.edu 80
Trying 128.238.26.21...
Connected to cis.poly.edu.
Escape character is '^]'.
GET /~ross/ HTTP/1.1
Host: cis.poly.edu
```

An HTTP request  
message

```
HTTP/1.1 200 OK
Date: Fri, 09 Aug 2019 02:46:08 GMT
Server: Apache/2.4.6
Last-Modified: Mon, 12 Nov 2018 16:25:17 GMT
ETag: "cf-57a7a257df256"
Accept-Ranges: bytes
Content-Length: 207
Content-Type: text/html; charset=UTF-8

<head>
<meta http-equiv="Content-Type" content="text/html; charset=windows-1252">
<meta http-equiv="refresh" content="0;url=http://nyu.edu/projects/keithwross/">
<title> Automatic Forwarding </title>
</head>
Connection closed by foreign host.
```

An HTTP response  
message, sent by the  
server.



```
ross.html x
1 <head>
2 <meta http-equiv="Content-Type" content="text/html; charset=windows-1252">
3 <meta http-equiv="refresh" content="0;url=http://nyu.edu/projects/keithwross/">
4 <title> Automatic Forwarding </title>
5 </head>
6
```



## Keith W. Ross



**Dean of Engineering and Computer Science, NYU Shanghai  
& Leonard J. Shustek Professor of Computer Science, CSE Dept,**

Keith Ross is the Dean of Engineering and Computer Science at NYU Shanghai. He is the Leonard J. Shustek Chair Professor of Computer Science at NYU. He also holds an affiliation with the Department of Computer Science at the Courant Institute of Mathematical Sciences and the Center for Data Science at NYU.

Previously he was a professor at University of Pennsylvania (13 years) and the University of Michigan (5 years). He was the Department Head of the CSE Department at the University of Michigan and he joined NYU Shanghai in 2013. He received a Ph.D. in Computer Science from The University of Michigan.

# User-server state: cookies

Many major Web sites use cookies to *identify users*.

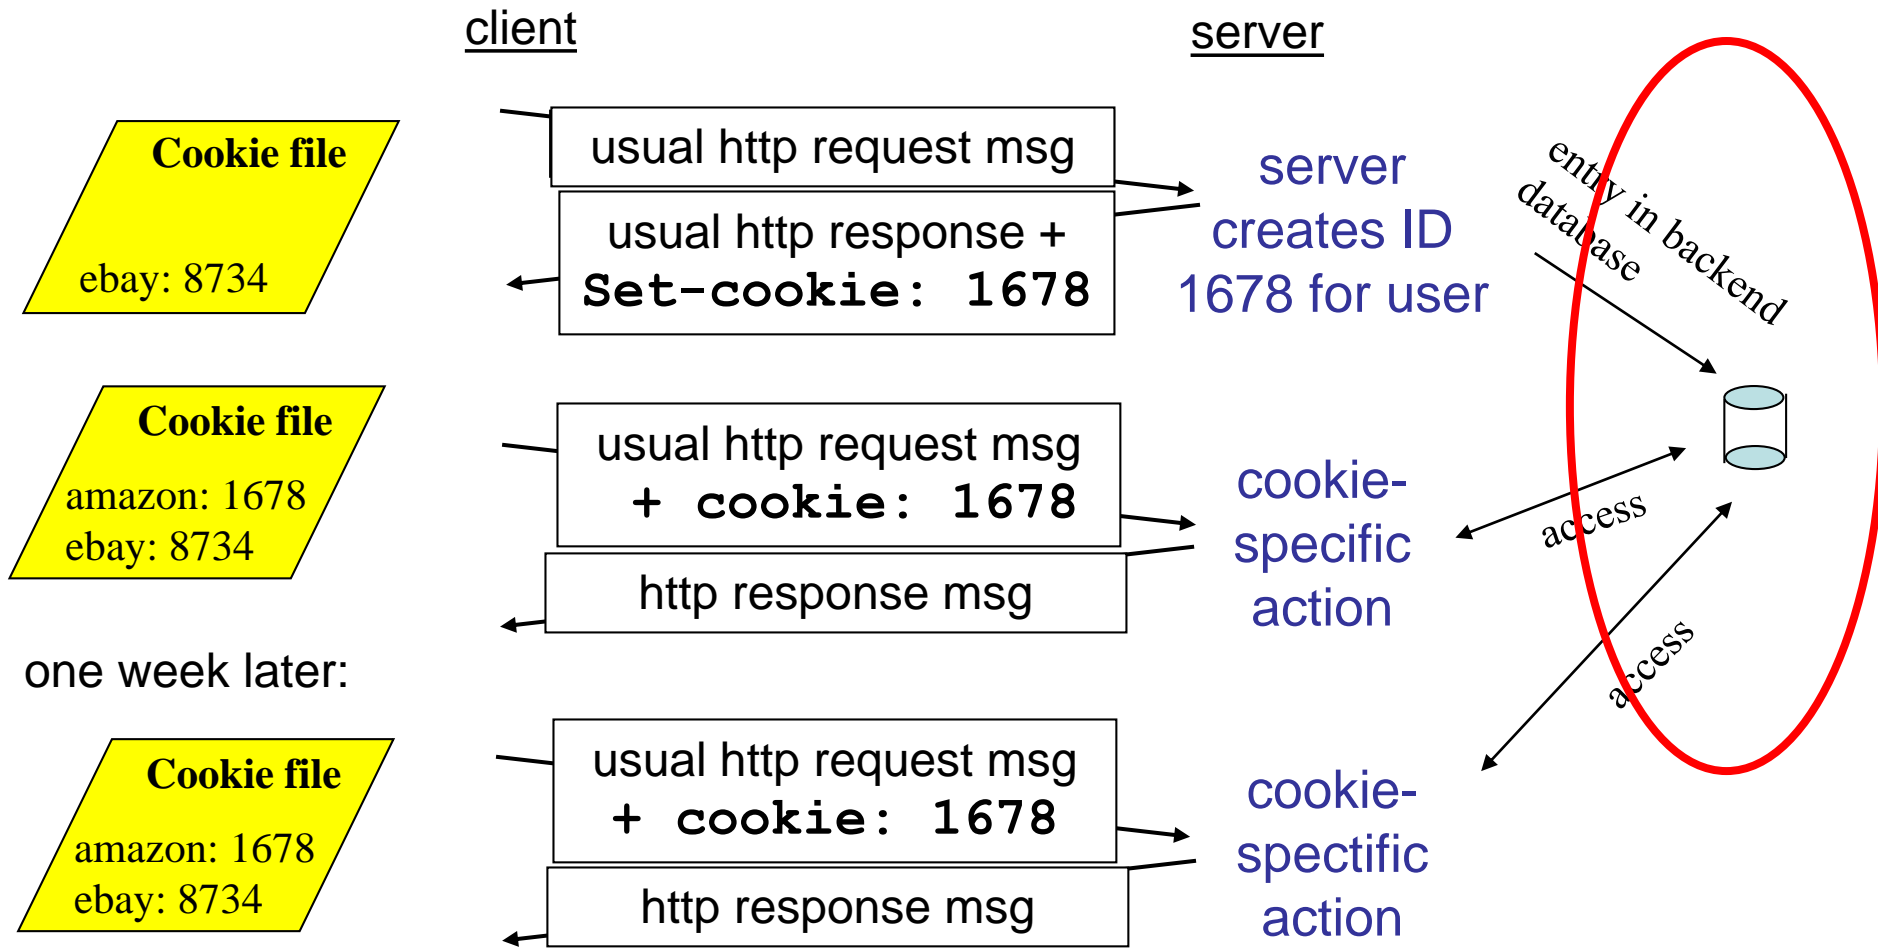
## Four components:

- 1) cookie header line in the HTTP response message
- 2) cookie header line in HTTP request message
- 3) cookie file kept on user's host and managed by user's browser
- 4) back-end database at Web site

### Example:

- Susan access Internet always from same PC
- She visits a specific e-commerce site for first time
- When initial HTTP requests arrives at site, site creates a unique ID and creates an entry in backend database for ID

# Cookies: keeping “state” (cont.)



# A cookie sample

← acm.org locally stored data	Remove All
__cfduid	^ ×
Name	__cfduid
Content	dee99f0a4dc0042015847bae526cfe9031564011334
Domain	.acm.org
Path	/
Send for	Any kind of connection
Accessible to script	No (HttpOnly)
Created	Thursday, July 25, 2019 at 11:35:36 AM
Expires	Friday, July 24, 2020 at 11:35:35 AM

# Cookies (continued)

## What cookies can bring:

- shopping carts
- recommendations
- user session state (Web e-mail)

— aside —


## Cookies and privacy:

1. cookies permit sites to learn a lot about you
2. you may supply name and e-mail to sites
3. search engines use redirection & cookies to learn yet more
4. advertising companies obtain info across sites

# Web caches (proxy server)

**Goal:** satisfy client request without involving origin server

- user sets browser: Web accesses via cache
- browser sends all HTTP requests to cache
  - object in cache: cache returns object
  - else cache requests object from origin server, then returns object to client




# More about Web caching

- Cache acts as both client and server
- Typically cache is installed by ISP (university, company, residential ISP)

## Why Web caching?

- Reduce response time for client request.
- Reduce traffic on an institution's access link.

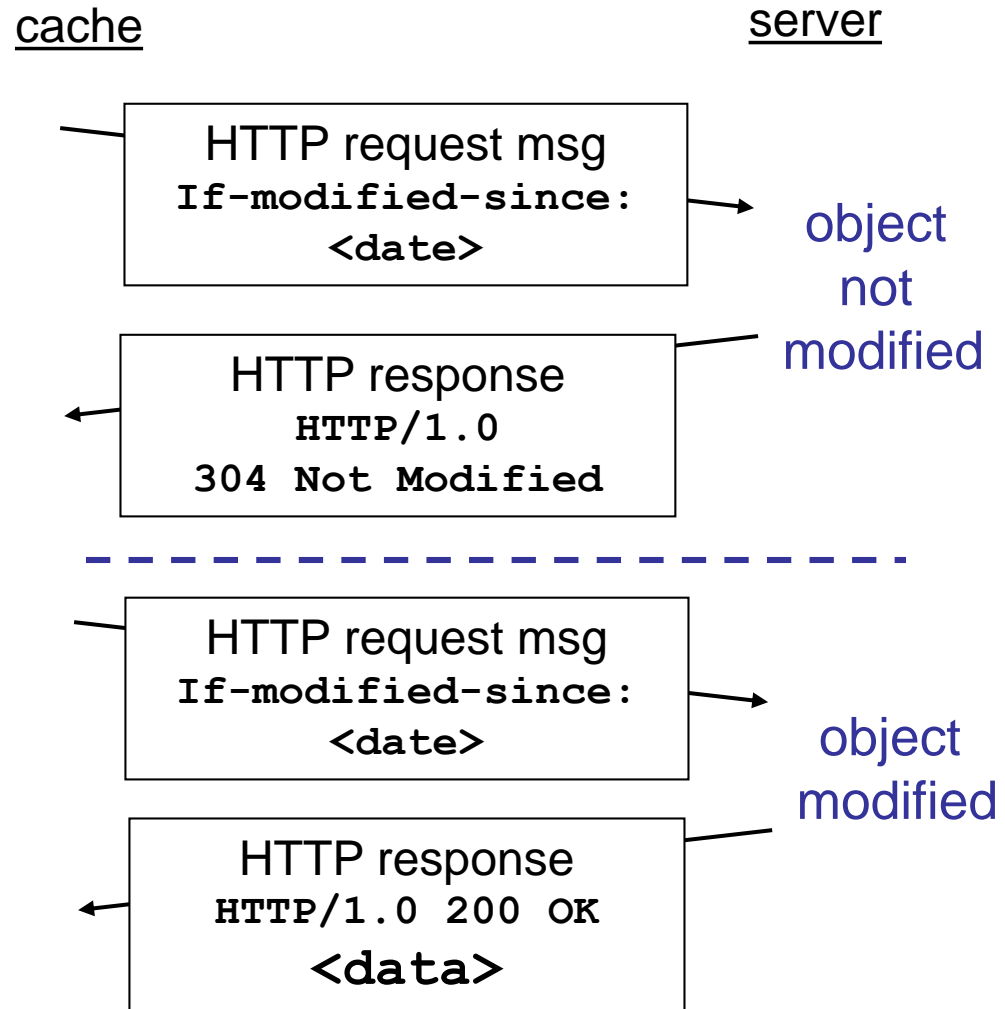


There is always new problem!  
The copy of an object in the cache may be stale.



# Conditional GET

- **Goal:** don't send object if cache has up-to-date cached version
- cache: specify date of cached copy in HTTP request  
`If-modified-since:`  
`<date>`
- server: response contains no object if cached copy is up-to-date:  
`HTTP/1.0 304 Not Modified`



# Summary

- Network applications
  - Network apps vs app. protocols
  - Application structure
- The Web
- HTTP

# References

- [KR3] James F. Kurose, Keith W. Ross, *Computer networking: a top-down approach featuring the Internet*, 3<sup>rd</sup> edition.
- [PD5] Larry L. Peterson, Bruce S. Davie, *Computer networks: a systems approach*, 5<sup>th</sup> edition
- [TW5] Andrew S. Tanenbaum, David J. Wetherall, *Computer network*, 5<sup>th</sup> edition
- [LHBi]Y-D. Lin, R-H. Hwang, F. Baker, *Computer network: an open source approach*, International edition

# Acknowledgements

- All slides are developed based on slides from the following two sources:
  - Dr DongSeong Kim's slides for COSC264, University of Canterbury;
  - Prof Aleksandar Kuzmanovic's lecture notes for CS340, Northwestern University,  
[https://users.cs.northwestern.edu/~akuzma/classes/CS340-w05/lecture\\_notes.htm](https://users.cs.northwestern.edu/~akuzma/classes/CS340-w05/lecture_notes.htm)